

UNIVERSITÀ DI PISA  
DIPARTIMENTO DI INGEGNERIA  
DELL'INFORMAZIONE

Master's Degree in Artificial Intelligence and Data  
Engineering

**BoardVerse**  
Large Scale and Multistructured Databases Project

Group:

**Martina Fabiani**  
**Tommaso Falaschi**  
**Emanuele Respino**

Project GitHub Repository:  
<https://github.com/martiFabia/BoardVerse/tree/Neo4j-Integration>

---

ACADEMIC YEAR 2024/2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
<b>2</b>	<b>Actors and Mockups</b>	<b>6</b>
2.1	Actors . . . . .	6
2.2	Mockups . . . . .	7
2.2.1	Unregistered User . . . . .	7
2.2.2	Registered User and Administrator . . . . .	7
<b>3</b>	<b>Requirements</b>	<b>16</b>
3.1	Functional Requirements . . . . .	16
3.2	Non-Functional Requirements . . . . .	19
<b>4</b>	<b>UML Class Diagram</b>	<b>20</b>
<b>5</b>	<b>Database</b>	<b>21</b>
5.1	Operations Volumes . . . . .	21
5.2	Document DB . . . . .	22
5.2.1	Collections . . . . .	22
5.2.2	Examples of documents . . . . .	23
5.2.3	CRUD Operations . . . . .	25
5.2.4	Queries . . . . .	27
5.3	Graph DB . . . . .	27
5.3.1	Nodes . . . . .	28
5.3.2	Relationship . . . . .	28
5.3.3	Examples of Graph . . . . .	29
5.3.4	CRUD Operations . . . . .	30
5.3.5	Queries . . . . .	32

---

<b>6 Distributed Database Design</b>	<b>33</b>
6.1 Replicas . . . . .	33
6.2 Handling Inter-DB Consistency . . . . .	35
6.3 Sharding . . . . .	36
<b>7 Dataset description</b>	<b>39</b>
7.1 Data retrieval process . . . . .	39
7.1.1 BoardGameGeek . . . . .	40
7.1.2 BoardGameArena . . . . .	40
7.2 Data processing . . . . .	41
7.3 Results . . . . .	41
<b>8 Implementation</b>	<b>43</b>
8.1 Spring Boot . . . . .	43
8.2 Security . . . . .	43
8.3 Config . . . . .	44
8.4 Model . . . . .	45
8.4.1 MongoDB models . . . . .	46
8.4.1.1 User Model . . . . .	46
8.4.1.2 Game Model . . . . .	47
8.4.1.3 Post Model . . . . .	47
8.4.1.4 Review Model . . . . .	48
8.4.1.5 Thread Model . . . . .	48
8.4.1.6 Tournament Model . . . . .	49
8.4.2 Neo4j models . . . . .	49
8.4.2.1 User Model . . . . .	49
8.4.2.2 Game Model . . . . .	50
8.4.2.3 Tournament Model . . . . .	51
8.5 Repository . . . . .	51
8.6 Service . . . . .	52
8.6.1 Analytics Service . . . . .	52
8.6.2 Auth Service . . . . .	52
8.6.3 Game Service . . . . .	52
8.6.4 Review Service . . . . .	53
8.6.5 Suggestion Service . . . . .	53
8.6.6 Thread Service . . . . .	53
8.6.7 Tournament Service . . . . .	53
8.6.8 User Service . . . . .	54

8.7	Controller . . . . .	54
8.7.1	Admin Controller . . . . .	54
8.7.2	Auth Controller . . . . .	54
8.7.3	Game Controller . . . . .	54
8.7.4	Review Controller . . . . .	54
8.7.5	Suggestion Controller . . . . .	54
8.7.6	Thread Controller . . . . .	55
8.7.7	Tournament Controller . . . . .	55
8.7.8	User Controller . . . . .	55
8.8	Exception . . . . .	55
8.9	Endpoints . . . . .	56
<b>9</b>	<b>Queries</b> . . . . .	<b>59</b>
9.1	MongoDB . . . . .	59
9.1.1	User operations . . . . .	59
9.1.2	Admin operations . . . . .	63
9.2	Neo4j . . . . .	66
9.2.1	Suggest new friends . . . . .	66
9.2.2	Suggest users with similar tastes . . . . .	67
9.2.3	Board games recommendation . . . . .	67
9.2.4	Tournaments recommendation . . . . .	68
9.3	Indexes . . . . .	69
9.3.1	MongoDB indexes . . . . .	69
9.3.2	Neo4j indexes . . . . .	72

# Chapter 1

## Introduction

BoardVerse is an application for true board game fans that allows users to discover, explore and share this passion. Thanks to a wide catalogue of games, users can search for new gaming experiences, read detailed reviews and ratings from other players, or share their opinions to help the community.

You can create a custom list of your favorite games, follow the latest trends, organize tournaments and connect with other fans with similar tastes. The platform also offers spaces for group discussions, where players can exchange strategies, tips and ideas.

This rich catalog is maintained by administrators, who not only add, update or remove games, but also supervise all user activities. For example, they can moderate inappropriate reviews or comments, ensuring that the app remains a welcoming environment and in line with the friendly spirit of the community.

The application uses a combination of technologies such as MongoDB for flexible management of documents and Neo4j to represent relationships between users, games and tournaments in an advanced way.

# Chapter 2

## Actors and Mockups

### 2.1 Actors

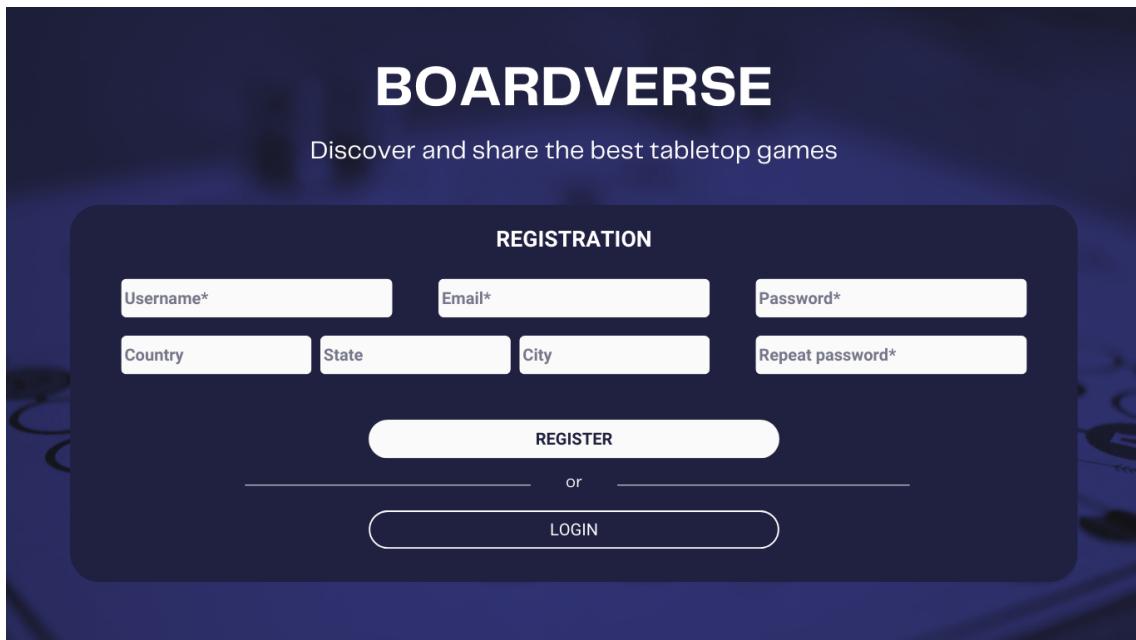
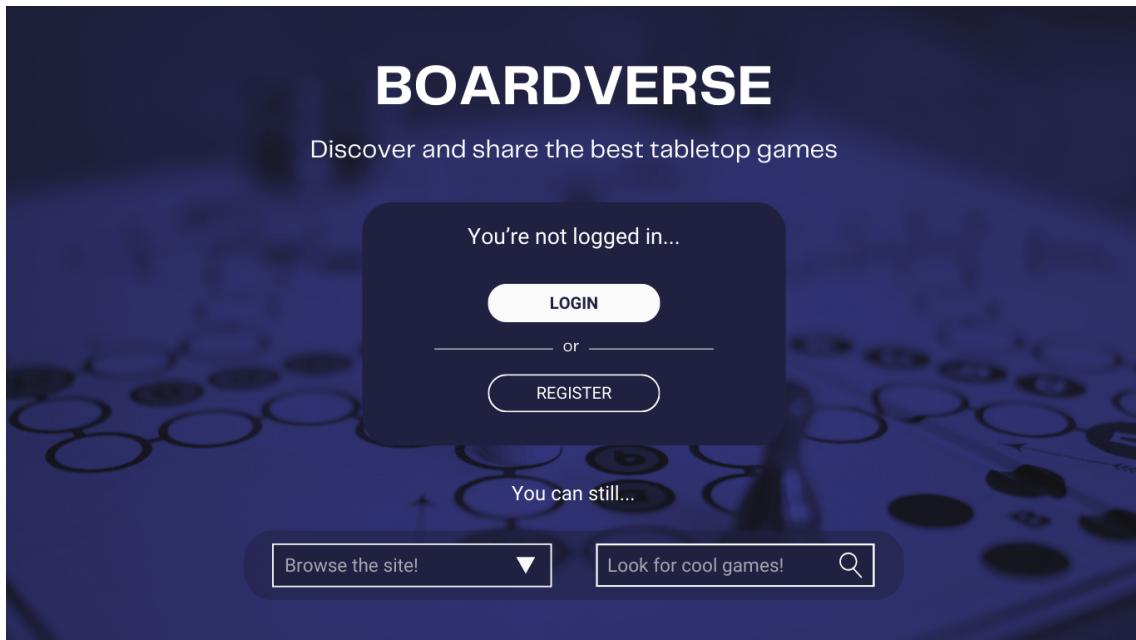
In BoardVerse there are three main actors:

- **Unregistered User**
- **Registered User**
- **Administrator**

The unregistered user has a limited set of actions that he can perform compared to the registered one who has access to all features of the application except those reserved for the administrator who will also manage the catalog of games and moderate the activity of registered users.

## 2.2 Mockups

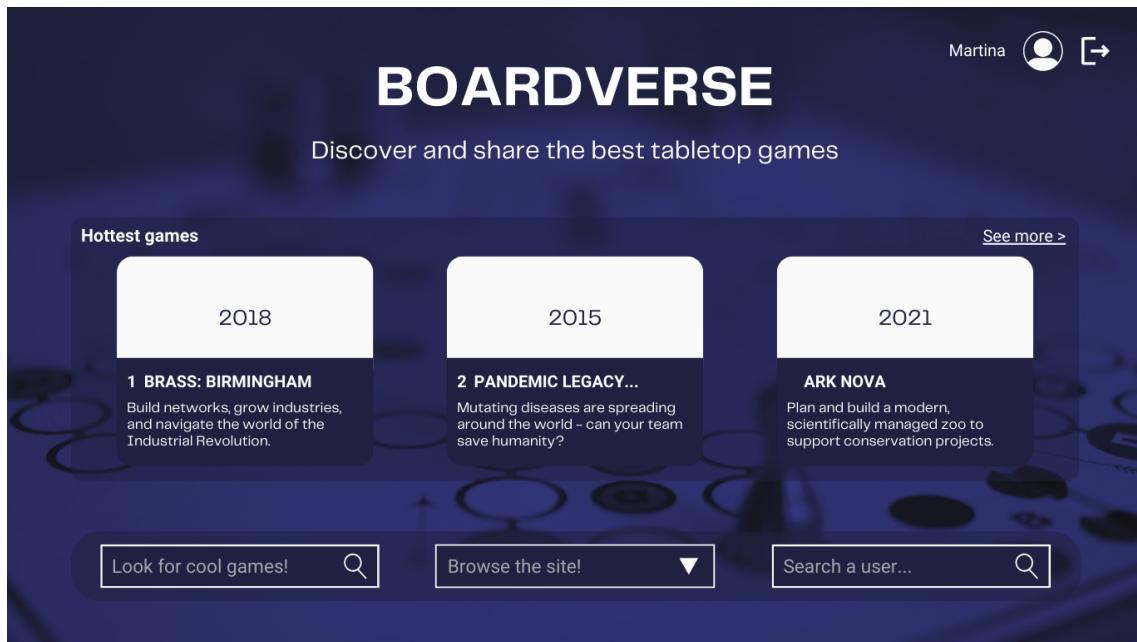
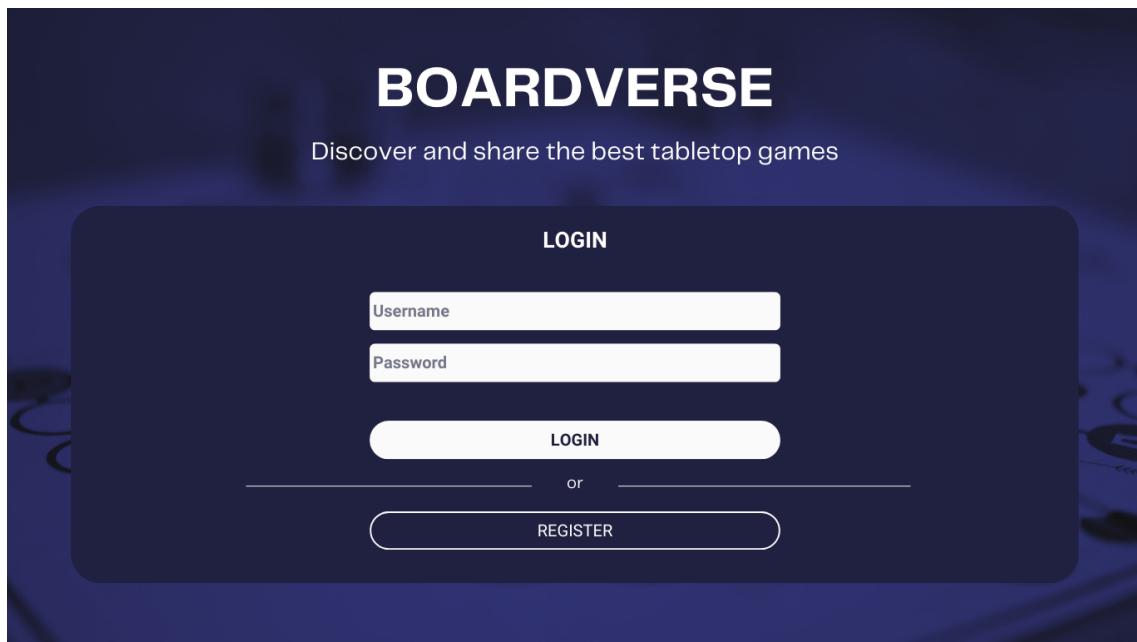
### 2.2.1 Unregistered User

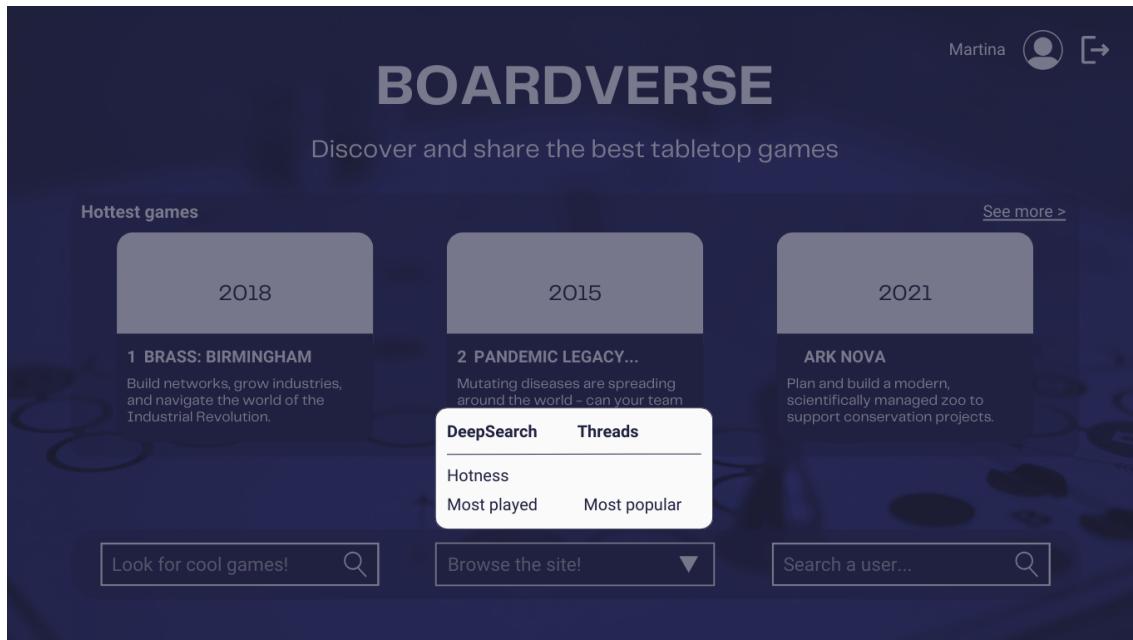


### 2.2.2 Registered User and Administrator

Actions that can be performed by both registered user and admin are highlighted with a red button, while actions that can only be done by

the admin are indicated with a green button. This is the main difference between the actions that can be carried out by the two different actors.





The page title "MOST POPULAR GAMES" is centered at the top. Navigation includes a "1-50 of" counter with arrows, and dropdown filters for "Country", "State or province", "City", "From: ---/---/---", and "To: ---/---/---".

Rank	Title	Avg Rating
1	<b>BANG! (2002)</b> Be the last one standing in this wild west shootout game with hidden roles.	9,9
2	<b>Pandemic Legacy: Season 1 (2015)</b> Mutating diseases are spreading around the world – can your team save humanity?	8,57
3	<b>Ark Nova (2021)</b> Plan and build a modern, scientifically managed zoo to support conservation projects.	8,53
4	<b>Gloomhaven (2017)</b> Vanquish monsters with strategic cardplay. Fulfill your quest to leave your legacy!	8,40

BOARDVERSE Browse the site! ▾ Look for cool games! 🔎 Martina  [→]

## DEEPSERACH

◀ ▶ 1-50 of Descending Ascending Year of release ▾ Category ▾ Mechanic ▾ SortBy ▾

Rank	Title	Avg Rating
1	<b>BANG! (2002)</b> Be the last one standing in this wild west shootout game with hidden roles.	9,9
2	<b>Pandemic Legacy: Season 1 (2015)</b> Mutating diseases are spreading around the world – can your team save humanity?	8,57
3	<b>Ark Nova (2021)</b> Plan and build a modern, scientifically managed zoo to support conservation projects.	8,53
4	<b>Gloomhaven (2017)</b> Vanquish monsters with strategic cardplay. Fulfill your quest to leave your legacy!	8,40

BOARDVERSE Browse the site! ▾ Look for cool games! 🔎 Martina  [→]

## THREADS

All Sessions General Rules News Reviews Strategy Variants Find players ◀ ▶ 1-50 Ascending Descending Game: ----- SortBy ▾ Year of release ▾ From: ---/---/--- To: ---/---/---

New Thread

- TERRAFORMING MARS - What's the best starting corporation for beginners?** Created: 2024-03-01 Last activity: 2025-01-05  
4 replies by MarsMaster
- CATAN - How to secure longest road every time?** Created: 2024-02-10 Last activity: 2024-12-05  
1 reply by RoadBuilderRick
- DIXIT - How to keep the clues challenging but not too obscure?** Created: 2024-03-01 Last activity: 2024-12-01  
2 replies by Imaginatively
- TICKET TO RIDE - Is it worth blocking routes to slow opponents down?** Created: 2024-02-12 Last activity: 2024-02-12  
0 replies by AndreTheMikeTysonFabbrino

BOARDVERSE

Browse the site! ▾

Look for cool games! 🔎

Martina 🌐 [→]

## THREADS

All >  
Sessions >  
General >  
Rules >  
News >  
Reviews >  
Strategy >  
Variants >  
Find players >

**Game: Bang!**

1-4 of 4

**Best strategy for the Sheriff role?** PARTECIPATE

Created: 2024-03-01  
Last activity: 2025-01-05

3 Strategy SheriffJoe

The Renegade can act friendly, but they're just waiting to backstab you! Published: 2024-03-01

- RenegadeRita REPLY

The Renegade can act friendly, but they're just waiting to backstab you! Published: 2024-03-02

- RenegadeRita REPLY

Agreed! Renegades are the wild cards of the game.

- OutlawSam REPLY

BOARDVERSE

Browse the site! ▾

Look for cool games! 🔎

Martina 🌐 [→]

 **BANG!** (2018) Avg Rating 9.9 ADD TO COLLECTION EDIT

Be the last one standing in this wild west shootout game with hidden roles.

**OVERVIEW** REVIEWS THREADS TOURNAMENTS

**Designer:** Emiliano Sciarra  
**Artist:** Stefano de Fazi, Pereyra il Tucumano  
**Publisher:** Arclight Games, Unipi...

**Categories:** American West, Bluffing...  
**Mechanisms:** Hand management, Hot potato...  
**Family:** Category: Dized Tutorial, Digital Implementations: Board Game Arena...

Playing time **20-40 min** Players **4-7** Age **8+**

Likes More statistics Rating details

**Description**  
*"The Outlaws hunt the Sheriff. The Sheriff hunts the Outlaws. The Renegade plots secretly, ready to take one side or the other. Bullets fly. Who among the gunmen is a Deputy, ready to sacrifice himself for the Sheriff? And who is a merciless Outlaw, willing to kill him? If you want to find out, just draw (your cards)!" (From back of box)*

The card game BANG! recreates an old-fashioned spaghetti western shoot-out... [see more >](#)

**BOARDVERSE**

Browse the site! ▾

Look for cool games! 🔎

Martina 🌐 [→]

 **BANG!** (2018)  
Avg Rating 9.9

Be the last one standing in this wild west shootout game with hidden roles.

OVERVIEW REVIEWS THREADS TOURNAMENTS

◀ ▶ 1-50 of 35.180 ADD REVIEW Date Rating

**martina13** ★★★★★ EDIT 10 - 2024/12/05  
Bang! is, without a doubt, the single most exhilarating and entertaining experience I've ever had in a board game. From the moment you draw your secret role, the game transforms into a dynamic blend of strategy, bluffing, and quick decision-making. The tension between...

**pohjanpalol17** ★★★★★ 9.7 - 2024/12/03  
No words needed.

**bangoa7** ★★★★★ 9.9 - 2024/11/20

**gabbiadini** ★★★★★ 9.8 - 2023/12/25

**BOARDVERSE**

Browse the site! ▾

Look for cool games! 🔎

Martina 🌐 [→]

 **BANG!** (2018)  
Avg Rating 9.9

Be the last one standing in this wild west shootout game with hidden roles.

OVERVIEW REVIEWS THREADS TOURNAMENTS

All > Sessions > General > Rules > News > Reviews > Strategy > Variants > Find players >

◀ ▶ 1-50 of 911 Ascending Descending SortBy ▾

**Bang! is the most amazing game ever!** Created: 2024-03-01 Last activity: 2025-01-05  
4 General theFactosAnnouncer

**Is the Dynamite card worth the risk?** Created: 2024-02-10 Last activity: 2024-12-05  
1 Rules lebonwski

**Best ways to spot the Renegade?** Created: 2024-03-01 Last activity: 2024-12-01  
2 General elBarto

**Top expansion pack for beginners?** Created: 2024-02-12 Last activity: 2024-02-12  
0 News luvumbo\_zito

**NEW THREAD**

**BOARDVERSE**

Browse the site! ▾ Look for cool games! 🔎 Martina 🚙

 **BANG!** (2018)  
Avg Rating: 9.9

Be the last one standing in this wild west shootout game with hidden roles.

OVERVIEW REVIEWS THREADS TOURNAMENTS

All Sessions General Rules News Reviews Strategy Variants Find players NEW THREAD

1-4 of 4 Game: Bang!

**Best strategy for the Sheriff role?** PARTECIPATE Created: 2024-03-01  
Strategy SheriffJoe Last activity: 2025-01-05

The Renegade can act friendly, but they're just waiting to backstab you! Published: 2024-03-01  
- RenegadeRita REPLY

The Renegade can act friendly, but they're just waiting to backstab you! Published: 2024-03-02  
- RenegadeRita

Aqreed! Renegades are the wild cards of the game.

**BOARDVERSE**

Browse the site! ▾ Look for cool games! 🔎 Martina 🚙

 **BANG!** (2018)  
Avg Rating: 9.9

Be the last one standing in this wild west shootout game with hidden roles.

OVERVIEW REVIEWS THREADS TOURNAMENTS NEW TOURNAMENT

1-50 of 2.840

**Bang! Showdown** Type: ELIMINATION Participants: 12/32 Start: 2024-03-05, 18:00 Organizer: Martina13 Open for registration

**Bang! Dodge City Special** Type: GROUP STAGE Participants: 16/16 Start: 2024-01-15, 15:30 Organizer: SheriffTina In Progress

**Bang! Beginner Brawl** Type: ROUND ROBIN Participants: 10/10 Start: 2024-03-05, 18:00 Organizer: DeputyDan Completed

The screenshot shows the BOARDVERSE website interface for a tournament titled "Bang! Showdown". At the top, there are navigation links for "Browse the site!" and "Look for cool games!", a user profile for "Martina" with a sign-in icon, and buttons for "REGISTER" and "MANAGE". Below the title, it says "Game: Bang! (2007)" and "Open for registration".

Key tournament details shown include:

- Start Date:** 2024-03-05, 18:00
- Winner:** -
- Organizer:** WildWestWill
- Visibility:** PUBLIC
- Type:** ELIMINATION
- Participants:** 12/32

The "Description" section contains the following information:

- Tournament Type:**
  - Single Elimination: Players face off in 1v1 matches. Winners advance to the next round until only one remains.
- Match Rules:**
  - Best of 3 games per match.
- Tie-breakers:**
  - In case of a tie, the player with the most damage dealt wins.

On the right side of the description section, there are three buttons: "Participants list", "Difficulty index", and "Social Density".

The screenshot shows the BOARDVERSE website interface for a user profile titled "PROFILE". The profile belongs to "Martina" and has an "ADMIN" status. At the top, there are navigation links for "Browse the site!" and "Look for cool games!", a user profile for "Martina" with a sign-in icon, and a button for "VIEW ANALYTICS".

The "USER DETAILS" section contains the following information:

Username:	martina13	EDIT
Email:	martina@example.com	
Registration date:	2024-01-15	
Birthday:	2001-06-14	
Location:	Pisa, Toscana, Italy	
Name:	Martina	
Surname:	Fabianski	

The "NETWORK" section shows:

- 15 following
- Maybe you know... (with a link to "Discovery similar users")
- 49 followers

The "RECENT REVIEWS" section lists three reviews for the game "BANG!" with a rating of 5 stars and the date 10 - 2024/12/5. Each review snippet ends with "see more >".

The "TOURNAMENTS" section shows:

- 9 participated
- 2 created
- 4 wins
- Suggest tournaments (with a link)

BOARDVERSE

Browse the site! ▾

Look for cool games! 🔎

Martina

VIEW ANALYTICS

**USER DETAILS**

Username: martina13

Email: martina@example.com

Registration date: 2024-01-15

Birthday: 2001-06-14

Location: Pisa, Toscana, Italy

Name: Martina

Surname: Fabianski

**RECENT REVIEWS**

**BANG!** ★★★★★ EDIT 10 - 2024/12/5

Bang! is, without a doubt, the single most exhilarating and entertaining experience I've ever had in a board game. From the moment you...

**BANG!** ★★★★★ EDIT 10 - 2024/12/5

Bang! is, without a doubt, the single most exhilarating and entertaining experience I've ever had in a board game. From the moment you...

**BANG!** ★★★★★ EDIT 10 - 2024/12/5

Bang! is, without a doubt, the single most exhilarating and entertaining experience I've ever had in a board game. From the moment you...

**NETWORK**

15 following

49 followers

Maybe you know...

Discovery similar users

**TOURNAMENTS**

9 participated 2 created 4 wins

*Suggest tournaments*

**Community geo-distribution**  
**Community tastes per age**  
**Monthly registrations**

[Game collection](#) [Discovery new games!](#)

# Chapter 3

# Requirements

## 3.1 Functional Requirements

- **Unregistered User**

1. The system must allow an unregistered user to browse games by name, category, year released and mechanics.
2. The system must allow an unregistered user to see the hottest, most popular, most played game rankings.
3. The system must allow an unregistered user to view detailed information about a selected game, including related reviews and threads.
4. The system must allow an unregistered user to browse threads by game, last post date and tag.
5. The system must allow an unregistered user to display messages in a specific thread.
6. The system must allow an unregistered user to register by providing personal information.

- **Registered User**

In addition to the previous requirements:

1. The system must allow a registered user to login by providing username and password.
2. The system must allow a registered user to edit personal information.

3. The system must allow a registered user to view their personal page.
4. The system must allow a registered user to find other registered user by username.
5. The system must allow a registered user to add a review to a game. A review can include a vote from 1 to 10, a comment or both.
6. The system must allow a registered user to edit a own review.
7. The system must allow a registered user to delete a own review.
8. The system must allow a registered user to create a thread for a game. A thread must include a subject, a tag and a content.
9. The system must allow a registered user to edit a own thread.
10. The system must allow a registered user to delete a own thread.
11. The system must allow a registered user to add a message or reply to a message in a thread.
12. The system must allow a registered user to delete a own message in a thread.
13. The system must allow a registered user to edit a own tournament.
14. The system must allow a registered user to delete a own tournament.
15. The system must allow a registered user to display tournaments related to a specific game for which he's allowed to.
16. The system must allow a registered user to view detailed information about a selected tournament, including participants, social density index and difficulty level.
17. The system must allow a registered user to create a tournament for a game. A tournament must include a name, type, description, location, minimum and maximum number of participants, start date and visibility.
18. The system must allow a registered user to participate to a tournament.

19. The system must allow a registered user to remove participation from a tournament.
20. The system must allow a registered user to update information related to a tournament of its own.
21. The system must allow a registered user to select a winner for a tournament of its own.
22. The system must allow a registered user to add a game to its own collection.
23. The system must allow a registered user to remove a game to its own collection.
24. The system must allow a registered user to follow another registered user.
25. The system must allow a registered user to unfollow another registered user.
26. The system must allow a registered user to see the list of followed users.
27. The system must allow a registered user to see the list of followers.
28. The system must allow a registered user to see the list of collected games.
29. The system must suggest a registered user a list of registered user who might know.
30. The system must suggest a registered user a list of registered user with similar tastes.
31. The system must suggest a registered user a list of games that may be of interest for the user.
32. The system must suggest a registered user a list of tournaments that may be of interest for the user.

- **Administrator**

In addition to all the previous requirements:

1. The system must allow ad administrator to add a game.
2. The system must allow ad administrator to update a game.

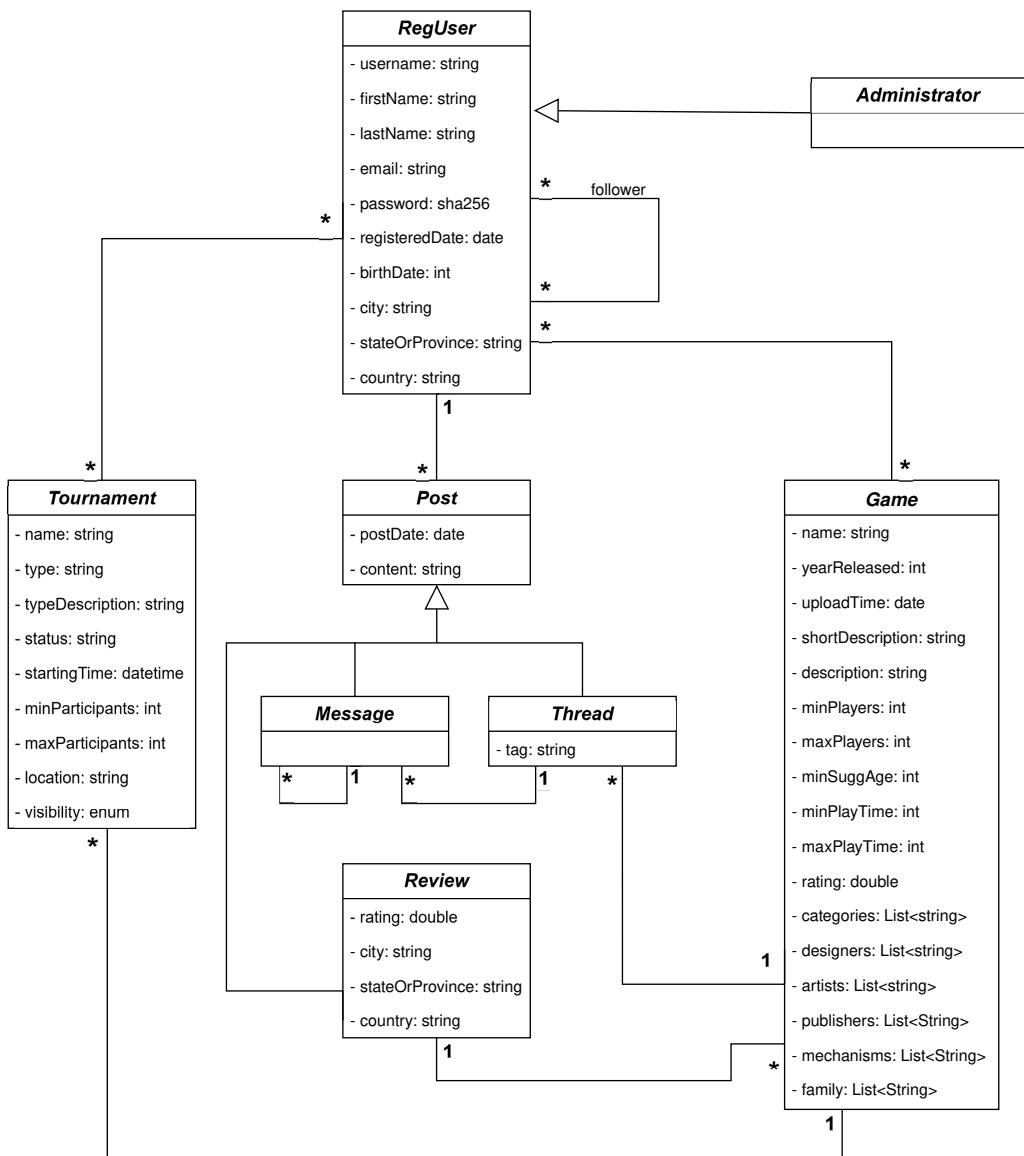
3. The system must allow ad administrator to delete a game.
4. The system must allow ad administrator to delete a user.
5. The system must allow ad administrator to delete a review.
6. The system must allow ad administrator to delete a thread.
7. The system must allow ad administrator to delete a message.
8. The system must allow ad administrator to delete a tournament.
9. The system must allow ad administrator to view the community geographical distribution analytic.
10. The system must allow ad administrator to view the community tastes per age bucket analytic.
11. The system must allow ad administrator to view the overall monthly registration distribution analytic.

## 3.2 Non-Functional Requirements

1. The system must be accessible as a web application compatible with major modern web browsers (Chrome, Firefox, Safari, Edge).
2. The system must be developed in an OOP language.
3. The system must ensure high availability.
4. The system should be capable of handling growth in user traffic by scaling effectively.
5. The system must be fault-tolerant and capable of maintaining essential functionality in the event of partial system failures.
6. The system must use a basic authentication protocol to secure access.
7. The system must encrypt users' passwords.

# Chapter 4

## UML Class Diagram



# Chapter 5

# Database

## 5.1 Operations Volumes

The following table presents an estimated number of occurrences for each of the primary operations performed by the application. These figures are based on an assumption of 100,000 active users per day.

Action	%	#Daily occurrences	MongoDB	Neo4j	Daily total
Register user	0.01	10	1	1	20
Login user	100	100,000	1	0	100,000
Browse users	20	20,000	1	0	20,000
User page load	80	$2 \times 80,000$	1	0	80,000
Update user info	0.01	10	4	1	40
Browse games	120	$3 \times 120,000$	1	0	360,000
Load game page	120	$3 \times 120,000$	1	0	360,000
Load ranking	50	50,000	1	0	50,000
Load hottest games	30	30,000	1	0	30,000
Write a review	10	10,000	3	0	30,000
Load game reviews	50	50,000	1	0	50,000
Create a thread	1	1,000	1	0	1,000
Write a message	5	5,000	1	0	5,000
Delete a message	0.01	10	1	0	10
Load game threads	30	30,000	1	0	30,000
Find threads	10	10,000	1	0	10,000
Load thread	25	$3 \times 25,000$	1	0	75,000
Create a tournament	0.1	100	2	1	300
Load game tournaments	25	25,000	1	0	25,000
Load tournament info	10	10,000	1	0	10,000
Follow a user	5	5,000	1	1	10,000
Unfollow a user	0.5	500	1	1	1,000
Game recommendation	20	20,000	0	1	20,000
User recommendation	15	15,000	0	1	15,000
Tournament recommendation	5	5,000	0	1	5,000
Like a game	10	10,000	0	1	10,000
Participate in a tournament	0.5	500	1	1	500

## 5.2 Document DB

### 5.2.1 Collections

- **USER:** Contains all documents related to registered users and admin with their personal information provided at the time of registration along with other useful fields for application purposes such as the following:
  - **registeredDate:** timestamp of registration.
  - **location:** country, state and city of the user.
  - **followers:** number of users following the user.
  - **following:** number of users followed by the user.
  - **tournaments:** number of created, participated and won tournaments.
  - **mostRecentReviews:** array of the three last reviews left by the user, made using the embedding of documents.
- **GAME:** The collection contains all documents about games information, with additional information among which we can find:
  - **uploadTime:** timestamp where the admin added the game.
  - **averageRating:** average rating of the game.
  - **ratingVoters:** number of votes given to the game.
- **REVIEW:** This collection contains all the reviews made for each game along with user metadata, such as:
  - **authorBirthDate:** Date of birth of the review author.
  - **location:** country, state and city of the user at the time of review.
- **THREAD:** collection containing thread-related documents. Threads contains the messages and information about the related game.
- **TOURNAMENT:** collection contains information about tournaments. One of its key fields is visibility, which determines who can access and participate in the tournament. The *visibility* field can have three possible values:

- PUBLIC: The tournament is open to everyone, and any user can visualize in tournaments list and join it.
- INVITE: Only users who receive a direct invitation from the administrator can participate.
- PRIVATE: The tournament is restricted, and only specific users can access it. The administrator explicitly lists the allowed users by storing their IDs in the *allowed* field.

### 5.2.2 Examples of documents

#### User document

```
_id: "40acdbcf-be6d-4e9c-a6bc-a23e3515d3fd"
username : "respino"
firstName : "Emanuele"
lastName : "Respino"
email : "emanuele.respino.4517@hushmail.com"
password : "87b555c0457fb41ebe56d877761376315014441693aece27297664aa02541aeb"
birthDate : 1956-05-01T00:00:00.000+00:00
location : Object
  city : null
  stateOrProvince : "Toscana"
  country : "Italy"
followers : 9
following : 40
tournaments : Object
  partecipated : 20
  won : 0
  created : 4
mostRecentReviews : Array (3)
  0: Object
    postDate : 2022-11-25T00:00:00.000+00:00
    rating : 7
    content : "Play with a friend who owns a copy. Brilliant game, terrific story and..."
    _id : "bed14528-79a1-44f1-950c-637b1950cbb0"
    game : Object
  1: Object
  2: Object
role : "ROLE_ADMIN"
registeredDate : 2024-06-22T00:00:00.000+00:00
```

## Game document

```
_id: "55cb996c-47bf-4b88-9e07-26a428ca6134"
name : "Squad Leader"
yearReleased : 1977
uploadTime : 2004-11-21T00:00:00.000+00:00
description: "A shot disturbs the eerie silence of a deserted city street, punctuate..."
shortDescription : "Command squads of soldiers in tactical WWII combat."
averageRating : 7.669209011299435
ratingVoters : 354
minPlayers : 2
maxPlayers : 2
minSuggAge : 14
minPlaytime : 60
maxPlaytime : 60
designers : Array (1)
  0: "John Hill (I)"
artists : Array (3)
  0: "Richard Hamblen"
  1: "Rodger B. MacGowan"
  2: "W. Scott Moores"
publishers : Array (3)
  0: "Arsenal Publishing, Inc."
  1: "The Avalon Hill Game Co"
  2: "Wydawnictwo Gołębiewski"
categories : Array (2)
  0: "Wargame"
  1: "World War II"
mechanics : Array (12)
family : Array (6)
```

## Review document

```
_id: "63a3a9ed-d61c-4e86-a5c2-6111d405b9cf"
postDate : 2022-09-14T00:00:00.000+00:00
location : Object
  city : null
  stateOrProvince : "Toscana"
  country : "Italy"
game : Object
  name : "Frosthaven"
  yearReleased : 2022
  shortDescription : "Adventure in the frozen north and build up your outpost throughout an ..."
  _id : "5f0df14a-8760-47b8-8130-8da8591af0c9"
  rating : 10
  content : "Lots of game play packed into this box."
  authorUsername : "respino"
  authorBirthDate : 1956-05-01T00:00:00.000+00:00
```

## Thread document

```

_id: "5198e8bd-6a5f-4bb1-9af7-352f780b3c69"
content: "Mini-review after playing tons of Lancashire ("In Your Face" -element ..."
postDate: 2024-09-05T19:27:02.000+00:00
lastPostDate: 2025-01-10T13:57:07.000+00:00
tag: "Reviews"
game: Object
messages: Array (11)
  0: Object
    postDate: 2024-09-05T19:27:02.000+00:00
    content: "Both are great games. As most people know, 80% of the rules are the sa..."
    authorUsername: "seepieceeggshell"
    _id: "74384a9e-74f6-456d-8bac-b542b8f591b3"
  1: Object
    postDate: 2024-09-06T00:04:38.000+00:00
    content: "I wouldn't say one game is "better" than the other because they are bo..."
    authorUsername: "keso55"
    _id: "a2043ad5-97ee-45a5-ac5e-ec222801afad"
  2: Object
  3: Object
  4: Object
  5: Object
  6: Object
  7: Object
  8: Object
  9: Object
  10: Object
authorUsername: "seepieceeggshell"

```

## Tournament document

```

_id: "3a294b18-2477-4110-b28b-42f531363976"
name: "The Round • Stephenson's Rocket month"
game: Object
  name: "Ark Nova"
  yearReleased: 2021
  _id: "b8e4e97d-e348-4ee4-a6d9-81590a4cbca2"
type: "Groups Stage"
typeDescription: "Players are first divided in groups for a first stage, winners advance..."
startingTime: 2025-01-26T16:00:00.000+00:00
location: Object
  city: null
  stateOrProvince: "California"
  country: "United States"
numParticipants: 225
minParticipants: 45
maxParticipants: 225
administrator: "davido"
winner: null
visibility: "PUBLIC"
options: Array (18)
  0: Object
    optionName: "Number of players in a match"
    optionValue: "2"
  1: Object
    optionName: "Number of players in a match (minimum)"
    optionValue: "2"
  2: Object
    optionName: "Mode of stage 1"
    optionValue: "Round robin"

```

### 5.2.3 CRUD Operations

- Create

- User.

- Game.
- Review.
- Thread.
- Tournament.

- **Read**

- User by id.
- User by username.
- Game by id.
- Game by name.
- Game by category.
- Game by yearReleased.
- Game by mechanic.
- Review by game id.
- Review by author username.
- Thread by id.
- Thread by game id.
- Thread by game yearReleased.
- Thread by postDate.
- Thread by tag.
- Tournament by id.
- Tournament by game id.

- **Update**

- Update user information.
- Update game information.
- Update review details.
- Update thread details.
- Add a new message in a thread.
- Reply to a message in a thread.

- Delete a message in a thread.
- Edit a message in a thread.
- Update tournaments details.

- **Delete**

- User.
- Game.
- Review.
- Thread.
- Tournament.

#### 5.2.4 Queries

Queries are described in more detail in the appropriate chapter.

- **Games ranking based on reviews filtered by postDate and user Location**
- **Top 10 Most Played Games**
- **Hottest games based on threads activity**
- **Game rating details**
- **Community geographical distribution**
- **Monthly user registrations by Year**
- **Community tastes per age bucket**

### 5.3 Graph DB

Within the graph we decided to track part of the activity of the user within the platform, starting from the interaction with other users through the follow feature, interaction with games allowing to place a game among the favorites and also everything that concerns the competitive side of a user with tournaments. Thanks to the implementation of the graph database it is possible to provide an improved and customized user experience.

### 5.3.1 Nodes

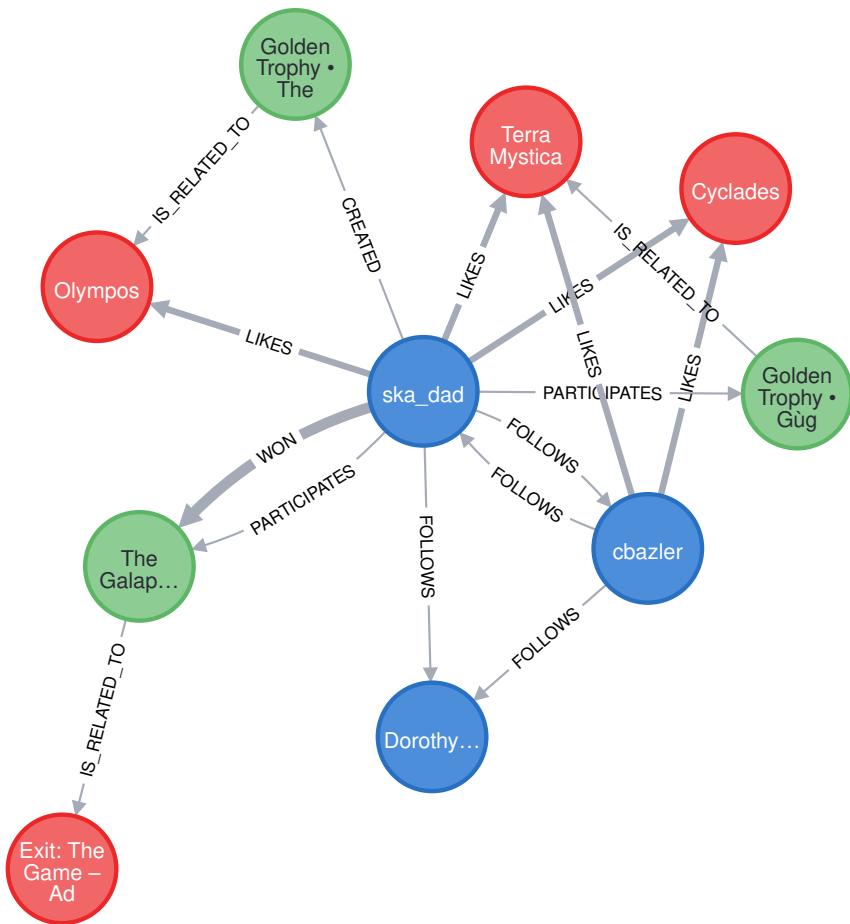
- **User:** Simple node for the registered user only contains the user-name.
- **Game:** node containing id, name, yearReleased, shortDescription, categories.
- **Tournament:** containing id, name, visibility, maxParticipants, startingTime.

### 5.3.2 Relationship

- **USER - FOLLOW → USER:** Relationship between two users created when one user starts following another one.
- **USER - LIKES → GAME:** Relationship between a user and a game, represents the inclusion of that game in the user's favorites list.
- **USER - CREATED → TOURNAMENT:** Tournament user relationship created when the user creates a tournament. The user who has this relationship with the tournament is the administrator of the tournament.
- **USER - PARTECIPATES → TOURNAMENT:** Relationship that specifies the user's participation in the tournament.
- **USER - WON → TOURNAMENT:** Relationship created at the end of a tournament that specifies a user's win in a tournament.
- **TOURNAMENT - IS RELATED TO → GAME:** Relationship created when a user creates a tournament that allows you to associate a particular tournament with the game it was created for.

All these relationships, except for IS RELATED TO, have the timestamp of when they were created.

### 5.3.3 Examples of Graph



### 5.3.4 CRUD Operations

<b>CREATE</b>	
<b>Domain-specific</b>	<b>Graph-centric</b>
Add a new User	Add new User vertex
Admin add a new Game	Add new Game vertex
User create a new Tournament	User add new Tournament vertex and add new "CREATED" relationship and a relationship "IS_RELATED_TO" is created from the tournament node to the game node.
User follows another user	Add new "FOLLOWS" relationships between the two users from the first user to the second one
User likes a Game	Add new "LIKES" relationship between User and Game
User Join a Tournament	Add new "PARTECIPATES" relationship between User and Tournament
User won a Tournament	Add new "WON" relationship between User and Tournament

<b>READ</b>	
<b>Domain-specific</b>	<b>Graph-centric</b>
Find all friends of an User	How many “FOLLOWSS” egde are outgoing from a specific User vertex?
Find all the tournaments a user is participating in	How many “PARTECIPATES” egde are outgoing from a specific User vertex?
Find all tournaments a user has won	How many “WON” egde are outgoing from a specific User vertex?
Find all tournaments that a user is a creator of	How many “CREATED” egde are outgoing from a specific User vertex?
Find all the games that the user likes	How many “LIKES” egde are outgoing from a specific User vertex?

<b>UPDATE</b>	
<b>Domain-specific</b>	<b>Graph-centric</b>
User update username	Update username property in the user node.
Creator of a tournament update the tournament info	Update tournament node properties
Admin update game info	Update game node properties

<b>DELETE</b>	
<b>Domain-specific</b>	<b>Graph-centric</b>
User is deleted	Remove a User node, all the Tournament nodes linked with "CREATED" relationship and all others "FOLLOWS", "PARTECIPATES" and "LIKES" relationship.
Game is deleted	Remove a Game node, all tournament nodes connected via "IS_RELATED_TO" relationship, and all "LIKES", "CREATED", "PARTECIPATES", "WON", "IS_RELATED_TO" relationships.
User removes a like to a game	Remove a "LIKES" relationship between User and Game
User no longer participates in a tournament	Remove a "PARTECIPATES" relationship between User and Tournament.
Tournament owner cancels the tournament	Remove a "TOURNAMENT" node and all its "OWNS", "PARTECIPATES" and "IS_RELATED_TO" connections
User removes follow to another user	Remove "FOLLOWS" relationship from first user to second one.

### 5.3.5 Queries

Queries are described in more detail in the appropriate chapter.

- Suggest new friends
- Suggest users with similar tastes
- Board games recommendation
- Tournaments recommendation

# Chapter 6

# Distributed Database Design

As outlined in section 3.2, our system is designed to guarantee High Availability and Partition Tolerance, ensuring fast responses to user requests. In alignment with the CAP theorem, the application adopts an AP-oriented approach, prioritizing availability and partition tolerance while compromising on consistency by implementing eventual consistency. In this chapter, we will analyze in detail the strategies for data distribution and replication, which form the foundation of the system's design.

## 6.1 Replicas

To ensure high availability, we have configured a replica set using the virtual machines available for this project. Specifically, the cluster includes three replicas dedicated to the document database (MongoDB). The organization of the replica set is as follows:

- **Primary node:** The node with the IP address 10.1.1.20 acts as the primary server and is responsible for handling all incoming requests. It has a priority level of 5, making it the preferred leader in the replica set.
- **Secondary nodes:** Two additional nodes, with IP addresses 10.1.1.21 (priority 2) and 10.1.1.22 (priority 1), store synchronized copies of

the data from the primary server. These nodes ensure data redundancy and support failover mechanisms in case the primary node becomes unavailable.

This setup allows the system to continue functioning even if the primary node becomes unavailable. In such a case, an automatic election process ensures that one of the secondary nodes is promoted to primary, minimizing downtime and maintaining service continuity.

To enhance system reliability, we have configured the nodes with different priority levels. The node with the highest priority (priority 5) is most likely to become the primary node during an election, ensuring a stable and predictable leadership structure. The node with the lowest priority (priority 1) also hosts the Neo4j service, and this configuration minimizes its workload by making it a candidate for primary only as a last resort. This approach carefully balances the workload between MongoDB and Neo4j, preventing resource conflicts and maintaining the performance of both databases.

Although the replica set currently includes three replicas exclusively for MongoDB and only one replica for Neo4j, this configuration is designed with future improvements in mind. While Neo4j's single replica is sufficient for this phase of the project, additional replicas can be added later to enhance its fault tolerance and availability. For MongoDB, the existing three-node cluster already provides strong support for high availability and redundancy.

Overall, the replica set design provides a solid foundation for a distributed database system that is highly available, fault-tolerant, and prepared for future scalability.

In our MongoDB configuration, we introduced the following constraints in the URI connection string:

```
mongodb://10.1.1.20:27017,10.1.1.21:27017,10.1.1.22:27017/?  
replicaSet=lsmdb&w=majority&readPreference=nearest&retryWrites=true
```

- **readPreference=nearest:** This parameter directs read operations to the nearest node in terms of network latency. This improves performance by reducing the time required to fulfill read requests, particularly in distributed environments. It also helps distribute the

workload more evenly across the replica set, as secondary nodes can participate in serving read requests, reducing the pressure on the primary node. While this configuration may occasionally return slightly stale data due to the eventual consistency model of MongoDB, it is a reasonable tradeoff in scenarios where speed and availability are prioritized.

- **w=majority:** This parameter ensures write operations are acknowledged only after being committed to the majority of nodes in the replica set. In our setup, this means at least two out of three nodes must confirm the write operation. This configuration guarantees data durability, as the system can recover from a primary node failure without losing data, provided the write has been replicated to the majority. Additionally, it maintains a high level of consistency, ensuring that the most recent updates are preserved across the replica set. While this approach may slightly increase write latency due to the need for replication confirmation, the tradeoff is well worth it for the added reliability and fault tolerance.

Together, these configurations create a balance between performance and resilience. The system delivers faster read responses by leveraging the nearest available node while ensuring that data remains safe and consistent through majority write acknowledgments.

## 6.2 Handling Inter-DB Consistency

Having two different types of databases in the system introduces some redundant data, which creates the need to carefully manage operations that interact with both databases. These operations, such as user registration or game information updates, require maintaining consistency between the two. To achieve this, we use **retryable** Spring functions that ensure operations are executed reliably.

The process begins with attempting the operation on the Neo4j database, as this is the more fragile system due to having only a single replica node. If the operation on Neo4j is successful, it then proceeds to execute the corresponding operation on the MongoDB database. If both operations are completed successfully, the function returns a success response to

the user. However, if either of the two operations fails, the system automatically retries the entire process up to three times, with a fixed delay between retries. If all retries fail, the operation is marked as unsuccessful, and an error response is returned to the user.

It's important to note that this approach does not include a rollback mechanism for changes made to Neo4j if the MongoDB operation fails. This could potentially lead to temporary inconsistencies between the databases. However, the likelihood of this happening is very low, given that MongoDB operates with a replica set of three nodes, providing greater reliability.

In the event that MongoDB fails during an operation, the retry mechanism starts from the very beginning of the function, including reattempting the write operation on Neo4j. This is not problematic because the write function is designed to be idempotent, meaning that repeating the operation will not create duplicate or conflicting entries. This design ensures that operations remain reliable and consistent across both databases, even in the face of potential failures.

## 6.3 Sharding

In our project, we have identified the collections that will benefit from sharding: Users, Games, Reviews, Threads and Tournaments. The selection of shard keys and the corresponding sharding strategies has been carefully designed, taking into account factors such as data access patterns, the anticipated growth rate of each collection, and the need to maintain optimal query performance. This approach ensures efficient data distribution and scalability across all collections, tailored to their specific usage scenarios.

- **Users collection**

**Sharding key:** Hashed Username

We have chosen to use the hashed username as the shard key to ensure a more uniform distribution of data across shards, preventing imbalances and facilitating scalability. This approach is also highly effective given the frequent queries based on the username, as it allows the system to directly route the query to the appropriate shard. This significantly enhances performance and efficiency.

- **Games collection**

**Sharding key:** Hashed id

The hashed ID has been chosen as the sharding key due to the high frequency of queries that rely on the ID to identify games. Additionally, this approach ensures seamless scalability. When new shards are added, data can be automatically redistributed across the cluster, thanks to the hashing mechanism, maintaining an optimal balance. This design enhances the system's resilience and makes it well-equipped to handle a growing volume of data efficiently.

- **Reviews collection**

**Sharding key:** Hashed gameID

For the Reviews collection, hashed gameID was chosen as the sharding key for several reasons:

- *Uniform Data Distribution:* By hashing the gameID, the reviews are distributed evenly across all shards, preventing any single shard from becoming overloaded. This is particularly important given that some games may have significantly more reviews than others.
- *Query Optimization:* Many queries in the Reviews collection are likely to focus on retrieving reviews for a specific game. Using gameID as the sharding key allows these queries to be directed to the relevant shard directly, significantly improving query performance by avoiding inter-shard operations.
- *Scalability:* The hashed gameID ensures that as the number of reviews grows, new shards can be added seamlessly, and the data will be automatically redistributed in a balanced way. This makes the system capable of handling increased data volumes efficiently.
- *Avoiding Hotspots:* Games with a large number of reviews could create hotspots if the data were not evenly distributed. Hashing the gameID breaks any natural patterns in the distribution of data, ensuring that the load is balanced across the cluster.

- **Threads collection**

**Sharding key:** Hashed gameID

As with the Reviews collection, the Threads collection also uses the hashed gameID as the sharding key. This choice was made because gameID is a field present in all documents, making it a reliable and consistent key for sharding. Additionally, the decision is supported by the same reasons outlined above, including uniform data distribution, query optimization, and improved scalability.

- **Tournaments collection**

**Sharding key:** Hashed gameID

Since tournaments can only be accessed from the corresponding game's page, we have chosen to use hashed gameID as the sharding key. This ensures that all tournament data related to a specific game is stored within the same shard, optimizing query performance by allowing direct access to the relevant shard when retrieving tournament information. Additionally, hashing the gameID ensures an even distribution of tournament data across shards, maintaining a balanced load and supporting scalability as the number of games and tournaments grows. This approach aligns with the hierarchical relationship between games and tournaments, streamlining data access and system efficiency.

# Chapter 7

## Dataset description

The core of this application revolves around managing and analyzing a large-scale dataset in the context of board games. To achieve this, it was essential to retrieve a robust and relevant dataset that would serve as the foundation for our application. For this purpose, we relied on two well-known platforms in the board game industry: **BoardGameGeek (BGG)** and **BoardGameArena (BGA)**. These platforms provided the raw material for our project about games, users, reviews, forums, and tournaments.

Our objective was to construct a dataset capable of simulating real-world scenarios while maintaining consistency across all data points. For data that was unavailable, we supplemented the dataset by injecting carefully crafted artificial data to ensure completeness and realism.

We can decompose this overall process, made using *Jupyter notebooks* and *Python* in two distinct and serialized phases: *data retrieval* and *data processing*.

### 7.1 Data retrieval process

First step was to download data from the up cited platforms. This required different techniques and some precautions due to policies, the huge amount of data available and time constrains. The overall raw data volume scraped is around **950MB**.

### 7.1.1 BoardGameGeek

BoardGameGeek provides a structured XML API (API2), which allowed us to efficiently gather various types of data. This platform was our primary source for:

- **Games.** We selected the top *2,000 games* based on their popularity ranking. For each game, we retrieved its detailed metadata.
- **Reviews.** During game data retrieval, user reviews were also collected for the selected games, including ratings, comments, and the associated usernames. To maintain consistency, we limited the number of reviews to a subset of the most recent ones (*500 maximum*) for each downloaded game. For our specific purpose, we downloaded the reviews which contained a text comment.
- **Threads.** Each game has associated forums with threads and messages: we retrieved forum data by first fetching the list of forums for a game, followed by the threads within each forum, and finally the individual messages within the threads. Forum activity was included to capture user engagement and community interactions around specific games. Due to the huge amount of data available, we only selected a reduced number of forums, avoiding the retrieval of the ones with less number of messages, from the first *300 games* in the original ranking, resulting in *70 threads per game* and *40 messages per thread* on average.
- **Users.** During reviews and threads data retrieval, we collected the usernames we found. Then, we selected *5000 users* for which we download the available metadata, including their favorite games.

### 7.1.2 BoardGameArena

Unlike BGG, BGA lacks an official API, requiring us to use web scraping to collect data. Custom scripts were developed from scratch to extract relevant tournament information, accordingly to rate limits and platform policies. However, this resulted in a limited data retrieval, downloading metadata for approximately *200 games*, with *10 tournaments per game* on average.

## 7.2 Data processing

Data has been initially stored in JSON format, simplifying its successive management and manipulation. To achieve our goal, we had to enhance data quality with many techniques such as *locations cleaning*, *introducing data redundancy*, *thread extraction* from forums and *messages refactoring*, ratings correction and so on. In addition, we needed to provide missing data injecting generated metadata: time constrains, tournaments replication, user activity redistribution to the downloaded users and a **user network** creation based on user similarity in their activity (likes and post on threads).

## 7.3 Results

The final dataset volume, approximately **1.2GB** is provided by several JSON files, reflecting the entities and the specific database in which then data has been stored:

File	Size
games.json	5.4MB
reviews.json	672.8MB
threads.json	223.4MB
tournaments.json	61.5MB
users.json	14.3MB

Table 7.1: File sizes for MongoDB

File	Size
games.json	589kB
tournaments.json	7.4MB
users.json	1.2MB
likes.json	123.9MB
follows.json	31.2MB
partecipates.json	86.2MB
created.json	3.2MB
won.json	1.5MB
is_related_to.json	2.6MB

Table 7.2: File sizes for Neo4j

# Chapter 8

# Implementation

In this chapter, we will delve into the detailed implementation of the project using Spring Boot, a powerful and widely used framework for building robust and scalable applications. The aim is to provide a clear and structured explanation of the project's architecture, focusing on the organization of the codebase and the specific roles of each component.

The implementation is divided into several key packages, each designed to address a specific concern in the application's structure. This modular approach ensures clarity, maintainability, and scalability.

## 8.1 Spring Boot

In this project, Spring Boot plays a crucial role by streamlining the development process with its robust framework. It simplifies dependency management, reduces boilerplate code, and provides an embedded server for running the application. Additionally, Spring Boot seamlessly integrates with various tools and technologies, enabling the efficient creation of REST APIs and ensuring scalability and maintainability.

## 8.2 Security

In this system, we configure Spring Security to ensure robust application security by implementing a range of security features. The authentication system supports login, registration, and role-based authorization, leveraging JWT (JSON Web Token) for secure token-based authentication, Spring Security for core security management, and Spring Data

MongoDB for efficient data handling.

To achieve this, we use a carefully designed security configuration, such as the following method:

```
@Bean
public SecurityFilterChain filterChain(HttpSecurity http) throws Exception {
    http.csrf(CsrfConfigurer<HttpSecurity>.csrf -> csrf.disable())
        .exceptionHandling(ExceptionHandlingConfigurer<HttpSecurity>.exception -> exception.authenticationEntryPoint(unauthorizedHandler))
        .sessionManagement(SessionManagementConfigurer<HttpSecurity>.session -> session.sessionCreationPolicy(SessionCreationPolicy.STATELESS))
        .authorizeHttpRequests(AuthorizationManagerRequestMatcher auth -> auth
            // Consente l'accesso agli endpoint Swagger
            .requestMatchers("/v3/api-docs/**", "/swagger-ui/**", "/swagger-ui.html").permitAll()
            .requestMatchers("/api/admin/**").hasRole("ADMIN")
            .requestMatchers("/api/users/**", "/api/suggestions/**").authenticated()
            .anyRequest().permitAll()
        );
    http.authenticationProvider(authenticationProvider());
    http.addFilterBefore(authenticationJwtTokenFilter(), UsernamePasswordAuthenticationFilter.class);
    return http.build();
}
```

This function plays a pivotal role in managing security configurations. It disables CSRF protection for stateless applications, sets up exception handling to manage unauthorized access, and configures session management to operate in a stateless mode. It also defines role-based access control, permitting unrestricted access to endpoints while restricting administrative endpoints to users with the "ADMIN" role.

Moreover, the function integrates an authentication provider and a JWT filter for secure authentication. This comprehensive setup ensures that the application is protected against unauthorized access while remaining user-friendly and maintaining seamless integration with essential tools like Swagger for API documentation. For declarative authorization management, the `@PreAuthorize` annotation is used. This is particularly valuable for endpoints that require diverse authorization rules, as it keeps the code clean, readable, and centralizes control over access permissions.

## 8.3 Config

This package contains essential configuration classes for the application. Specifically:

- **GlobalExceptionHandler:** This class centralizes the management of exceptions across the application. It ensures consistent error

handling and structured responses for various types of exceptions, improving maintainability and user experience.

- **SwaggerConfig:** This class configures Swagger to generate and manage API documentation for the application. It simplifies the process to explore and test the API, while seamlessly integrating security features like JWT authentication into the documentation. This ensures a clear understanding of secured endpoints and provides an interactive interface for testing the API with proper authentication.

## 8.4 Model

The model package defines the data structures used throughout the application. The system's models are divided into two types:

- **MongoDB**
- **Neo4j**

## 8.4.1 MongoDB models

### 8.4.1.1 User Model

```
@Document(collection = "users") * Martina +1 *
@Data
public class UserMongo {
    @Id
    private String id;

    @NotBlank
    @Size(max = 20)
    @Indexed(unique = true)
    private String username;
    @NotBlank(message = "First name is required")
    private String firstName;
    @NotBlank(message = "Last name is required")
    private String lastName;
    @NotBlank
    @Size(max = 50)
    @Email(message = "Email should be valid")
    private String email;
    @NotBlank
    @Size(min = 8, max = 20)
    private String password;
    private Date registeredDate;
    @Past(message = "Birthdate must be in the past")
    private Date birthDate;
    @Schema(description = "UserMongo Location (Country, State, City)")
    private Location location;
    private int followers;
    private int following;
    private TournamentsUser tournaments;
    private Role role; // Ruolo enum: "ROLE_USER" o "ROLE_ADMIN"
    private List<ReviewUser> mostRecentReviews=new ArrayList<>();
```

### 8.4.1.2 Game Model

```
@Document(collection = "games") 29 usages ▲ LeBonWskii +1 *
@Data
public class GameMongo {

    @Id
    private String id;

    @NotBlank(message = "Name cannot be blank")
    @Indexed
    private String name;
    @NotNull(message = "Year released cannot be null")
    private Integer yearReleased;
    private Date uploadTime;

    private String description;
    private String shortDescription;
    @Min(value = 0, message = "Average rating must be positive")
    private Double averageRating;
    private int ratingVoters;
    private int minPlayers;
    private int maxPlayers;
    private int minSuggAge;
    private int minPlayTime;
    private int maxPlayTime;
    private List<String> designers;
    private List<String> artists;
    private List<String> publisher;
    @NotNull(message = "Categories cannot be null")
    private List<String> categories;
    private List<String> mechanics;
    private List<String> family;
```

### 8.4.1.3 Post Model

```
@Data 5 usages 3 inheritors ▲ Martina
public abstract class Post {
    private String authorUsername;
    private Date postDate;
    private String content;
}
```

#### 8.4.1.4 Review Model

```
@Document(collection = "reviews") 12 usages ▾ Martina +1
@Data
public class ReviewMongo extends Post {

    @Id
    private String id;

    private Date authorBirthDate;
    private Location location;
    private GameReview game;

    @Min(1) @Max(value = 10, message = "Rating must be between 1 and 10")
    @NotNull
    private Double rating;
```

#### 8.4.1.5 Thread Model

```
@Document(collection = "threads") 11 usages ▾ Martina +2
@Data
public class ThreadMongo extends Post {
    private String id;
    private String tag;
    private GamePreviewEssential game;
    private Date lastPostDate;

    private List<Message> messages;
}
```

### 8.4.1.6 Tournament Model

```
@Data 15 usages  ↗ Martina +1
@Document(collection = "tournaments")
public class TournamentMongo {
    @Id
    private String id;
    @NotBlank
    private String name;
    @NotBlank
    private GamePreviewEssential game;

    private String type;
    private String typeDescription;
    private Date startingTime;

    private Location location;

    @Positive
    private Integer numParticipants;
    @Positive
    private Integer minParticipants;
    @Positive
    private Integer maxParticipants;

    private String administrator;
    private String winner;

    private VisibilityTournament visibility;

    private List<OptionsTournament> options=new ArrayList<>();
    private List<String> allowed;
```

## 8.4.2 Neo4j models

### 8.4.2.1 User Model

```
✓ @Node("User") 9 usages  ↗ LeBonWskii
@Data
public class UserNeo4j {

    @Id
    @Property("username")
    private String username;
}
```

### 8.4.2.2 Game Model

```
@Node("Game") 13 usages  ↗ LeBonWskii +1
@Data
public class GameNeo4j {
    @Id
    @Property("_id")
    private String id;

    @Property("name")
    @NotBlank
    private String name;

    @Property("yearReleased")
    @NotNull
    @PositiveOrZero
    private int yearReleased;

    @Property("shortDescription")
    @NotBlank
    private String shortDescription;

    @Property("categories")
    @NotEmpty
    private List<String> categories;
}
```

#### 8.4.2.3 Tournament Model

```
@Node("Tournament") 3 usages  ↗ LeBonWskii +1
@Data
public class TournamentNeo4j {
    @Id
    @Property("_id")
    private String id;

    @Property("name")
    @NotBlank
    private String name;

    @Property("visibility")
    @NotBlank
    private String visibility;

    @Property("maxParticipants")
    @NotNull
    @Positive
    private int maxParticipants;

    @Property("startingTime")
    @NotNull
    private Date startingTime;
}
```

## 8.5 Repository

The repository modules in our system play a critical role in facilitating interaction with both MongoDB and Neo4j databases. Leveraging Spring Boot's repository features, they provide a abstraction layer enabling efficient implementation of CRUD functionalities and simplifying data management.

For MongoDB and Neo4j, we utilize annotations like `@Query` for custom queries and `@Aggregation` for handling advanced operations such as analytics and recommendations. These features ensure flexibility in accessing and processing data while maintaining high performance.

## 8.6 Service

The service layer is the backbone of our application, bridging the gap between the controllers and repositories. It serves as the central hub for business logic, validating data, managing transactions, and ensuring a clear separation of responsibilities.

Service modules are responsible for applying business rules, verifying data integrity, and handling potential errors before interacting with the repository layer. They act as a safeguard, ensuring that only valid and compliant operations are executed on the database.

Beyond simply coordinating database interactions, the service layer processes and transforms data retrieved from repositories, making it ready for use by the controller or other parts of the application. By centralizing key functionalities within the service layer, the system achieves greater modularity, maintainability, and resilience, enabling consistent and efficient data management while adhering to business requirements.

### 8.6.1 Analytics Service

Module to provide the administrator with analytics functionality for generating statistics for better understanding of application performance, For example, features that allow users to be displayed across countries and track the monthly trend of user registrations.

### 8.6.2 Auth Service

Module that implements the registration and login of a user that through encryption of passwords and JWT tokens allows to increase the security of the application in these delicate phases. Also if a user registration is successful, it creates a new document in the User collection in MongoDB and a new node in Neo4j.

### 8.6.3 Game Service

Manages all CRUD operations related to games, from creating a new game, to updating the information about it, deleting and searching for

games through numerous filters including name, release year and categories, returning lists of games with the support of pagination. In addition to this are implemented some queries such as obtaining the ranking of games based on rating, even in this case filterable for numerous parameters, and get the most played games.

#### 8.6.4 Review Service

Allows you to manage all the CRUD operations related to reviews, such as creation of new reviews, updating of the review by the author, Deletion of a review by the administrator's author or obtaining reviews for a specific game or user.

#### 8.6.5 Suggestion Service

It allows to generate important suggestions to improve the user experience on the platform by exploiting the connections of the graph database Neo4j. In particular, it allows a user to suggest other users whom he or she might know or have interests in communities or to suggest tournaments which the user might be interested in participating in.

#### 8.6.6 Thread Service

Implements all thread management, starting from the creation, deletion, update and addition of messages or replies allowing users to interact exchanging advice and opinions. It also allows you to get threads related to a game or to get threads simply by filtering through other parameters such as the date of creation of the thread to allow users to always stay up-to-date with the latest discussions.

#### 8.6.7 Tournament Service

Manages all the features related to tournaments both MongoDb side and NEo4j side, allowing you to perform CRUD operations related to tournaments that allow you to implement the competitive side of table games loved by users,

### 8.6.8 User Service

It realizes all those functionalities related to the users, such as the user's profile, view of a user's profile or personal profile allowing also to modify the information of the latter. In addition to the deletion of the profile it is also possible to obtain the references left by the user.

## 8.7 Controller

### 8.7.1 Admin Controller

Provides all the functionality needed by the admin, such as game management and moderation of user activity. It also allows the administrator to get important analytics on the application. This controller is connected with game, user and analytic services for a management of the key parts of the application.

### 8.7.2 Auth Controller

Manages the parts related to the authentication of a user by connecting to the auth service to ensure greater security of the login and registration phases.

### 8.7.3 Game Controller

Allows to manage all requests concerning the games by connecting with the game service to allow a correct return of what the user had requested.

### 8.7.4 Review Controller

Connects with the review service to provide the functionality of creating, obtaining, updating and deleting reviews.

### 8.7.5 Suggestion Controller

Provides useful suggestion functionality to users by leveraging the suggestion service.

### 8.7.6 Thread Controller

It allows to map all user requests regarding thread related features, such as creating, deleting or adding messages to populate the discussion. To do this, of course, uses the thread service.

### 8.7.7 Tournament Controller

Requests for creation, updating, deletion and obtaining, both for each game that a user, for tournaments are provided also to the connection with tournament service.

### 8.7.8 User Controller

This controller is responsible for handling various user-related functionalities, such as searching for users, fetching user details, updating their profiles, and deleting accounts. It also manages user preferred games, allowing users to add or remove a like to a certain game. Furthermore, it oversees social interactions like following or unfollowing other users. The controller communicates with the User Service to ensure efficient processing of all requests.

## 8.8 Exception

This package contains custom exception classes:

- **AlreadyExistsException:** This exception is likely thrown when an entity (e.g., a user or game) already exists in the system and a duplicate operation is attempted (e.g., trying to register a user with an existing username).
- **NotFoundException:** This exception is thrown when a requested entity or resource cannot be found (e.g., looking for a non-existent user or game).

## 8.9 Endpoints

### Authentication User authentication operations

POST	/api/auth/signup		
POST	/api/auth/login		

### Admin Admin operations

POST	/api/admin/games Add a new game		
DELETE	/api/admin/games/{id} Delete a game		
PATCH	/api/admin/games/{id} Update a game		
GET	/api/admin/analytics/usersByLocation Get number of users by country		
GET	/api/admin/analytics/monthlyRegistrations Get monthly registrations by year		
GET	/api/admin/analytics/bestGamesByAge Get best games by age		
DELETE	/api/admin/users/{username} Delete a user		

### User Operations related to user management

POST	/api/users/{username}/follow Start following a user		
DELETE	/api/users/{username}/follow Stop following a user		
GET	/api/users/myProfile Get user profile		
DELETE	/api/users/myProfile Delete user profile		
PATCH	/api/users/myProfile Update user profile		
GET	/api/users/{username} Get user info		
GET	/api/users/{username}/wonTournaments Get list of won tournaments by a user		
GET	/api/users/{username}/participatedTournaments Get list of participated tournaments by a user		
GET	/api/users/{username}/organizedTournaments Get list of organized tournaments by a user		
GET	/api/users/{username}/likedGames Get games liked by a user		
GET	/api/users/{username}/following Get users followed by a user		
GET	/api/users/{username}/followers Get followers of a user		
GET	/api/users/myProfile/reviews Get user reviews		
GET	/api/users/browse Browse users		

Game Operations related to games		
POST	/api/games/{gameId}/like Like a game	🔒 ↴
DELETE	/api/games/{gameId}/like Unlike a game	🔒 ↴
GET	/api/games/{gameId} Get game page	🔒 ↴
GET	/api/games/{gameId}/statistics Get games statistics about likes and tournaments	🔒 ↴
GET	/api/games/{gameId}/ratingsDetail Get game ratings detail	🔒 ↴
GET	/api/games/{gameId}/likes Get the list of user who liked a game	🔒 ↴
GET	/api/games/ranking Get games ranking	🔒 ↴
GET	/api/games/mostPlayed Get the best 10 most played games	🔒 ↴
GET	/api/games/hottest Hottest games based on threads activity	🔒 ↴
GET	/api/games/filter Find games by filter	🔒 ↴
GET	/api/games/browse Browse games	🔒 ↴

Review Operations related to reviews		
GET	/api/games/{gameId}/reviews Get game reviews	🔒 ↴
POST	/api/games/{gameId}/reviews Add a new review	🔒 ↴
DELETE	/api/games/{gameId}/reviews/{reviewId} Delete a review	🔒 ↴
PATCH	/api/games/{gameId}/reviews/{reviewId} Update a review	🔒 ↴

Suggestion Operations related to suggestions		
GET	/api/suggestions/suggestedUsers	🔒 ↴
GET	/api/suggestions/suggestedTournaments	🔒 ↴
GET	/api/suggestions/suggestedGames	🔒 ↴
GET	/api/suggestions/similarUsers	🔒 ↴

Thread Operations related to threads		
GET	/api/games/{gameId}/threads Get game threads	🔒 ↴
POST	/api/games/{gameId}/threads Add a new thread	🔒 ↴
PATCH	/api/threads/{threadId}/messages Add a new message to a thread	🔒 ↴
PATCH	/api/threads/{threadId}/messages/{ messageId }/reply Reply to a message	🔒 ↴
PATCH	/api/threads/{threadId}/messages/{ messageId }/edit Edit a message	🔒 ↴
PATCH	/api/threads/{threadId}/messages/{ messageId }/delete Delete a message	🔒 ↴
GET	/api/threads/{threadId} Get a thread	🔒 ↴
DELETE	/api/threads/{threadId} Delete a thread	🔒 ↴
GET	/api/threads/filter Find threads by filter	🔒 ↴

Tournament Operations related to tournaments		^
POST	/api/games/{gameId}/tournament Add a new tournament	🔒 ✓
POST	/api/games/tournaments/{tournamentId}/selectWinner Select a winner for a tournament	🔒 ✓
POST	/api/games/tournaments/{tournamentId}/register Register to a tournament	🔒 ✓
DELETE	/api/games/{gameId}/tournament/{tournamentId} Delete a tournament	🔒 ✓
PATCH	/api/games/{gameId}/tournament/{tournamentId} Update a tournament	🔒 ✓
GET	/api/games/{gameId}/tournaments Get game tournaments	🔒 ✓
GET	/api/games/tournaments/{tournamentId} Get tournament detail	🔒 ✓
GET	/api/games/tournaments/{tournamentId}/socialDensity Get tournament social density index	🔒 ✓
GET	/api/games/tournaments/{tournamentId}/participants Get tournament participants	🔒 ✓
GET	/api/games/tournaments/{tournamentId}/difficulty Get tournament difficulty index	🔒 ✓
DELETE	/api/games/tournaments/{tournamentId}/unregister Unregister from a tournament	🔒 ✓

# Chapter 9

## Queries

### 9.1 MongoDB

In the following chapter, we will illustrate the queries implemented in MongoDB, providing detailed explanations of their structure and functionality.

#### 9.1.1 User operations

##### **Ranking of games based on reviews filtered by postDate and user Location**

This query is useful for generating a ranked list of games based on user reviews within a specific time frame and location.

For example:

- Ranking games by average rating for users in a particular country or city.
- Generating a leaderboard of games based on reviews during a specific promotional period or event.

It provides insights into the most highly-rated games within a given context and time frame, helping identify popular games. Each entry in the result represents a game and includes the following information: game Id, game Name, Short Description, Year Released and average rating.

The aggregation performs the following steps:

1. Match: Filters reviews to include only those within the specified date range (startDate to endDate). Additional location-based filters are applied dynamically, ensuring that optional parameters such as country, state, and city are handled flexibly.
2. Group: Groups reviews by the unique game identifier (gameId). For each group, the aggregation retrieves the game information and calculates the average rating of the reviews for that game.
3. Sort: Orders the grouped results in descending order of the calculated averageRating, ensuring the highest-rated games appear first.

```

@Aggregation(pipeline = { 1usage new *
    "{$match": { "postDate": { "$gte": ?0, "$lte": ?1 } } },
    "{$match": { "$and": [
        + " { '$or': [ { '$expr': { '$eq': [:#country, null] } }, { 'location.country': :#{country} } ] },
        + " { '$or': [ { '$expr': { '$eq': [:#state, null] } }, { 'location.stateOrProvince': :#{state} } ] },
        + " { '$or': [ { '$expr': { '$eq': [:#city, null] } }, { 'location.city': :#{city} } ] }
    + " ] } },
    "{$group": {
        + " '_id': '$game._id', 'name': { '$first': 'game.name' },
        + " 'yearReleased': { '$first': 'game.yearReleased' },
        + " 'shortDescription': { '$first': 'game.shortDescription' },
        + " 'averageRating': { '$avg': 'rating' }
    + " ] } },
    "{$sort": { 'averageRating': -1 } },
    "{$project": {
        + " '_id': 1, 'name': 1, 'shortDescription': 1, 'yearReleased': 1,
        + " 'averageRating': { '$round': [ '$averageRating', 2 ] }
    + " } }"
})
Slice<GameRankPreviewDTO> findAverageRatingByPostDateLocation(
    Date startDate, Date endDate, String country, String state, String city, Pageable pageable);

```

## Top 10 Most Played Games

This aggregation query calculates and retrieves the Top 10 games with the highest average weighted participation based on tournaments held within a specified time frame. The query incorporates weighted calculations based on the visibility of the tournaments (e.g., public, invite, or private). Here's a detailed explanation of its steps:

1. Match: Filters the tournaments to include only those that started within the specified date range and excludes tournaments with fewer than 10 participants.
2. AddFields: Adds two new fields for each tournament:
  - Weight*: Assigns a weighting factor based on the tournament's visibility (PUBLIC, INVITE, or PRIVATE).

*WeightedParticipants*: Calculates the weighted number of participants by multiplying numParticipants by the corresponding weight.

3. Group: Groups the tournaments by the unique identifier of the associated game.
4. AddFields: Calculates the average weighted participation for each game by dividing the total weighted participants by the total weight.
5. Sort: Orders the games in descending order of their average weighted participation, ensuring the most played games appear first.
6. Limit: keep only the top 10

```

@Aggregation(pipeline = { 1 usage  ± Martina
    "{'$match': {'$and': [{ '$gt': { 'startingTime': ?0, 'gameID': ?1 }, 'numParticipants': { '$gt: 10 } } ]",
    "{'$addFields': {'weight': { '$switch': { 'branches': [ " +
        "{'$case': { '$eq': ['$visibility', 'PUBLIC'] }, 'then': 2 }, " +
        "{'$case': { '$eq': ['$visibility', 'INVITE'] }, 'then': 1.5 }, " +
        "{'$case': { '$eq': ['$visibility', 'PRIVATE'] }, 'then': 1 }, " +
        "'default': 1 } ], " +
        "'weightedParticipants': { '$multiply': ['$numParticipants', " +
        "{$switch': { 'branches': [ " +
            "{'$case': { '$eq': ['$visibility', 'PUBLIC'] }, 'then': 2 }, " +
            "{'$case': { '$eq': ['$visibility', 'INVITE'] }, 'then': 1.5 }, " +
            "{'$case': { '$eq': ['$visibility', 'PRIVATE'] }, 'then': 1 }, " +
            "'default': 1 } ] } } }",
    "{'$group': { '_id': '$game._id', 'name': { '$first': '$game.name' }, 'yearReleased': { '$first': '$game.yearReleased' }, " +
        "'totalWeightedParticipants': { '$sum': '$weightedParticipants' }, 'totalWeight': { '$sum': 'weight' } } }",
    "{'$addFields': { 'averageWeightedParticipants': { '$divide': ['$totalWeightedParticipants', '$totalWeight'] } } },
    "{'$sort': { 'averageWeightedParticipants': -1 } }",
    "{'$limit': 10 }",
    "{'$project': { '_id': 0, 'gameID': '_id', 'name': 1, 'yearReleased': 1, 'averageWeightedParticipants': 1 } }"
})
List<MostPlayedGameDTO> findTop10GamesWithHighestAverageParticipation(Date startDate, Date endDate);
}

```

## Hottest games based on threads activity

This query is designed to rank games based on the importance of associated messages, taking into account both the number of messages and their time relevance within a specified period. It provides a normalized view of the game's activity or significance, potentially useful for identifying trending games or games with the most impactful discussions during a specific time frame. The aggregation performs the following steps:

1. addFields: the query first filters the messages array within each document, retaining only messages where the postDate falls within the specified date range.

2. addFields: next, it processes the filtered messages array using map. For each message:
  - A variable dateDif is calculated, representing the difference between the range limits.
  - The weight of each message is computed using the formula:  

$$0.1 + (1 * ((\text{postDate} - \text{startDate}) / (\text{endDate} - \text{startDate})))$$
 which normalizes the weight based on the postDate relative to the reference range. The computed weight is then merged back into the original message object.
3. project: a new field, totalWeight, is added to each document. This field represents the sum of all calculated weights for the messages in the messages array.
4. group: the documents are grouped by gameId. For each group: the game details are retrieved and the totalWeight values for all messages associated with the game are summed to compute an importanceIndex.
5. project: in the final stage, only the required fields are selected
6. sort: the results are sorted in descending order of importanceIndex, ensuring that games with the highest importance appear first.

```
@Aggregation(pipeline = {
    @Usage(EmmanuelRsp)
    "{$addFields: { messages: { $filter: { input: '$Messages', as: 'message', cond: { $and: [ { $gte: ['$message.postDate', ?0] }, { $lte: ['$message.postDate', ?1] } ] } } } }",
    "{$addFields: { messages: { $map: { input: '$Messages', as: 'message', in: { $let: { vars: { maxDate: ?1, minDate: ?0, dateDif: { $subtract: [?1, ?0] } } },
        "in: { $mergeObjects: [ '$$message', { weight: { $add: [?1, { $divide: [ { $multiply: [?1, { $subtract: ['$message.postDate', ?0] } ] } ] } ] } } } } } } }",
    "{$project: { game: 1, totalWeight: { $sum: '$messages.weight' } } }",
    "{$group: { _id: '$game._id', game: { $first: '$game' }, importanceIndex: { $sum: '$totalWeight' } } }",
    "{$project: { _id: '_id', name: '$game.name', yearReleased: '$game.yearReleased', importanceIndex: 1 } }",
    "{$sort: { importanceIndex: -1 } }"
})
Slice<BestGameThreadDTO> getNormalizedGameRankingsWithDetails(Date startDate, Date endDate, Pageable pageable);
```

## Game rating details

This query is designed to calculate and retrieve rating statistics for a specific game based on its gameId. It provides a detailed breakdown of average ratings, the standard deviation of ratings, and a distribution of ratings by rounded values. Here's an explanation of the query stages:

1. match: filters documents to include only ratings associated with the specified gameId, ensuring the rating field is not null.
2. project: retains the original rating and adds a roundedRating field, which rounds the rating to the nearest whole number for distribution analysis.
3. facet: Splits into two parallel computations:
  - Stats: Calculates the average rating (avgRating) and standard deviation (stdDeviation) across all ratings.
  - Distribution: Groups ratings by roundedRating, counts occurrences in each category, and sorts them in ascending order.
4. project: Combines the stats and distribution results into a single document, providing both a statistical summary and a detailed distribution of ratings for the game.

```

@Aggregation(pipeline = { @usage ▾ Martina
    "{$match": { "game._id": ?0, "rating": { "$ne": null } } }",
    "{$project": { "rating": 1, "roundedRating": { "$round": [ "$rating", 0 ] } } }",
    "{$$facet": { " +
        "'stats': [ " +
            "{$group": { "_id": null, "avgRating": { "$avg": "$rating" }, "stdDeviation": { '$stdDevPop': '$rating' } } } " +
        "], " +
        "'distribution': [ " +
            "{$group": { "_id": '$roundedRating', 'count': { '$sum': 1 } } }, " +
            "{$sort": { '_id': 1 } }, " +
            "{$project": { 'rating': '_id', 'count': 1, '_id': 0 } } " +
        "] " +
        "}" },
    "{$$project": { " +
        "'avgRating': { '$arrayElemAt': ['$stats.avgRating', 0] }, " +
        "'stdDeviation': { '$arrayElemAt': ['$stats.stdDeviation', 0] }, " +
        "'distribution': '$distribution' " +
    "}" }"
)
RatingDetails findRatingDetailsByGameId(String gameId);

```

## 9.1.2 Admin operations

### Community geographical distribution

This query calculates the number of users grouped by country and state or province, organizing the results into a hierarchical structure with states nested under their respective countries. It is ideal for generating a location-based user distribution, providing insights into how users are geographically distributed by country and state/province. It can be

used for dashboards, reporting, or analytics to understand regional user engagement. The aggregation performs the following steps:

1. match: Filters the documents to include only users with valid location information.
2. group: groups the users by country and stateOrProvince, calculating the total number of users for each state or province.
3. group: regroups the results by country, creating a hierarchical structure where each country contains an array of its states along with their respective user counts.
4. sort: orders the results alphabetically by country for better readability and presentation.
5. project: formats the final output to include: country and states, an array of states within the country, each showing the user count.

```
@Aggregation(pipeline = { 1usage ± Martina
    "{$match": { "location.country": { "$ne": null }, "location.stateOrProvince": { "$ne": null } } },
    "{$group": { "_id": { "country": "$location.country", "state": "$location.stateOrProvince"}, "count": { "$sum": 1 } } },
    "{$group": { "_id": "$_id.country", "states": { "$push": { "state": "$_id.state", "count": "$count" } } } },
    "{$sort": { "_id": 1 } }", // Ordenamiento alfabetico su country
    "{$project": { "_id": 0, "country": "$_id", "states": 1 } }"
})
Slice<CountryAggregation> countUsersByLocationHierarchy(Pageable pageable);
```

## Monthly user registrations by Year

This query is designed for the admin and provides the monthly user registrations for a selected year. It aggregates the number of users registered per month within the specified time range, allowing the admin to analyze registration trends for the chosen year. This query is specifically tailored for admin users to:

- Visualize monthly registration trends for a given year.
- Identify months with high or low registration activity.
- Analyze the effectiveness of campaigns or events influencing user registrations during specific periods.

Here's how it works:

1. match: filters the user documents to include only those whose registeredDate falls within the specified year.
2. group: Groups the documents by month and calculates the total number of user registrations for each month.
3. sort: Orders the results in ascending order based on the month number.

The query produces a list of objects, where each object contains:

- Month: The numeric representation of the month (e.g., 1 for January, 2 for February) .
- Registrations: The total number of user registrations for that specific month.

```
@Aggregation(pipeline = { 1 usage  ± Martina
    "{$match": { "registeredDate": { "$gte": ?0, "$lt": ?1 } } }",
    "{$group": { "_id": { "$month": "$registeredDate" }, "registrations": { "$sum": 1 } } },
    "{$sort": { "_id": 1 } }"
})
List<MonthlyReg> monthlyRegistrations(Instant startDate, Instant endDate);
```

## Community tastes per age bucket

This query is designed to determine the best-rated games grouped by age brackets of users who provided ratings. It calculates the average rating for each game within defined age groups and selects the game with the highest average rating for each bracket. It helps developers and publishers tailor their content, marketing strategies, or updates to better suit specific audience segments. Below is an explanation of how the query works:

The query filters reviews to include only those with a non-null rating, ensuring that only valid data is processed. It then calculates a new field, age, by determining the user's age as the difference in years between their birth date (authorBirthDate) and the current date. Based on this, it creates an additional field, ageBracket, grouping users into 10-year age ranges. Reviews are then grouped by both age bracket and game ID,

and games with fewer than 30 reviews are excluded to ensure statistical significance. The results are sorted first by age bracket and then by average rating in descending order. For each age bracket, the game with the highest average rating is selected. Finally, the age brackets are sorted in ascending order to produce the final output.

```

@Aggregation(pipeline = {
    usage & Martina
    "{$match: { rating: { $ne: null } } }",
    "{$addFields: { age: { $dateDiff: { startDate: '\"$authorBirthDate\"', endDate: '\"$NOW\"', unit: '\"year\"' } } } }",
    "{$match: { age: { $gte: 10, $lte: 99 } } }",
    "{$addFields: { ageBracket: { $concat: [ \"",
        "{$toString: { $multiply: [ 10, { $floor: { $divide: [ '\"$age\", 10 ] } } ] } }, \"-\", \"",
        "{$toString: { $add: [ { $multiply: [ 10, { $floor: { $divide: [ '\"$age\", 10 ] } } ], 9 ] } } ] } } }",
    "{$group: { \" +",
        "_id: { ageBracket: '\"$ageBracket\"', game: '\"$game._id\"' }, \" +",
        "name: { $first: '\"$game.name\"' }, \" +",
        "yearReleased: { $first: '\"$game.yearReleased\"' }, \" +",
        "averageRating: { $avg: '\"$rating\"' }, \" +",
        "totalReviews: { $sum: 1 } \" +",
    "}}",
    "{$match: { totalReviews: { $gte: 30 } } }",
    "{$sort: { \"_id.ageBracket\": 1, \"averageRating\": -1 } }",
    "{$group: { \" +",
        "_id: '\"$_id.ageBracket\"', \" +",
        "gameID: { $first: '\"$_id.game\"' }, \" +",
        "name: { $first: '\"$name\"' }, \" +",
        "yearReleased: { $first: '\"$yearReleased\"' }, \" +",
        "bestAvgRating: { $first: '\"$averageRating\"' }, \" +",
        "totalReviews: { $first: '\"$totalReviews\"' } \" +",
    "}}",
    "{$sort: { \"_id\": 1 } }"
})
List<BestGameAgeDTO> findBestGameByAgeBrackets();

```

## 9.2 Neo4j

### 9.2.1 Suggest new friends

**Domain-specific:** Suggest new friends to user A, suggesting users who have more friends in common with user A but are not friends with user A

**Graph-centric:** Considering the current User vertex (User A), identify all the User vertices directly connected to User A through an outgoing "FOLLOWS" edge. Then, for each of these connected User vertices, find all other User vertices connected to them through an outgoing "FOLLOWS" edge. Count the occurrences of these User vertices to determine the number of mutual friends they share with User A. Exclude any User vertex that is already directly connected to User A via a "FOLLOWS"

edge or is User A itself. Sort the remaining User vertices by the number of mutual friends in descending order and return them as potential new friends.

```
@Query("""
    MATCH (currentUser:User {username: $username})-[:FOLLOWS]->(followedUser:User)-[:FOLLOWS]->(suggestedUser:User)
    WHERE NOT (currentUser)-[:FOLLOWS]->(suggestedUser) AND suggestedUser.username <> $username
    RETURN suggestedUser.username AS username, COUNT(suggestedUser) AS commonFollowers
    ORDER BY commonFollowers DESC
    SKIP $pageSize * ($pageNumber - 1)
    LIMIT $pageSize
""")
List<UserFollowRecommendationDTO> getUsersRecommendationBySimilarNetwork(String username, int pageSize, int pageNumber);
```

## 9.2.2 Suggest users with similar tastes

**Domain-specific:** Suggest users with similar tastes to user A, suggesting users who have more liked games in common with user A but are not friends with user A

**Graph-centric:** Considering the current User vertex (User A), identify all the User vertices directly connected to User A through an outgoing "FOLLOWS" edge. Then, for each of these connected User vertices, find all other User vertices connected to them through an outgoing "FOLLOWS" edge. Locate all the games liked by these users. Then locate all games connected with a "LIKES" relationship with user A. Count the games liked in common between user A and other users. Return the list of users who have at least 1 liked game in common with user A in descending order of liked games in common

```
@Query("""
    MATCH (user:User {username: $username})-[:FOLLOWS]->(:User)-[:FOLLOWS]->(suggestedUser:User)
    WHERE NOT (user)-[:FOLLOWS]->(suggestedUser) AND suggestedUser <> user
    MATCH (user)-[:LIKES]->(commonGame:Game)<-[:LIKES]->(suggestedUser)
    WITH suggestedUser, COUNT(commonGame) AS sharedGames
    WHERE sharedGames > 0
    RETURN
        suggestedUser.username AS username,
        sharedGames
    ORDER BY sharedGames DESC
    SKIP $pageSize * ($pageNumber - 1)
    LIMIT $pageSize
""")
List<UserTastesSuggestionDTO> getUsersRecommendationBySimilarTastes(@Param("username") String username, @Param("pageSize") int pageSize, @Param("pageNumber") int pageNumber);
```

## 9.2.3 Board games recommendation

**Domain-specific:** Board games recommendation based on the user's favorite board games and categories

**Graph-centric:** Starting from the current user node A, all other user nodes connected to the A node are identified through a "FOLLOWS" relationship. Between these nodes identify all games connected through

a "LIKES" relationship that are not connected with a "LIKES" relationship with the user node A. Then find all the game nodes connected to the A node through a "LIKES" relationship. Extract all categories of potentially suggestible games and games liked by the user. Compare and count the number of categories in common. Return games with at least 1 common category in descending order of the number of common categories.

```

@Query("""
    MATCH (u:User {username: $username})--[:FOLLOWS]-(follower:User)-[:LIKES]->(g:Game)
    WHERE NOT (u)-[:LIKES]->(g)
    WITH DISTINCT g, u
    // Compute common categories
    MATCH (u)-[:LIKES]->(likedGame:Game)
    UNWIND likedGame.categories AS userCategory
    UNWIND g.categories AS gameCategory
    WITH g, userCategory, gameCategory
    WHERE userCategory = gameCategory
    WITH g, COUNT(userCategory) AS commonCategories
    // Order and return results
    WHERE commonCategories > 0
    RETURN
        g.name AS name,
        g._id AS id,
        g.yearReleased AS yearReleased,
        g.shortDescription AS shortDescription,
        commonCategories
    ORDER BY commonCategories DESC
    SKIP $pageSize * ($pageNumber - 1)
    LIMIT $pageSize
""")
List<GameSuggestionDTO> getGamesRecommendation(String username, int pageSize, int pageNumber);

```

### 9.2.4 Tournaments recommendation

**Domain-specific:** Tournament recommendation based on the favorite games of the user and tournaments that the user's friends are participating in

**Graph-centric:** Starting from the user A vertex, search for all tournament nodes linked via a "PARTICIPATES" relationship and find all games linked to these tournaments via a "IS \_ RELATED \_ TO" relationship. All tournament summits linked to previously found games that have public visibility, are not yet started and are not full where user A does not participate already. Count the user vertices connected to user A through a "FOLLOWS" relationship that also have a "PARTICI-

PATES" relationship with the tournaments found. Return tournaments in descending order of the number of followers participating, and in case of tie give priority to tournaments that start earlier.

```

@Query("""
    MATCH (user:User {username: $username})-[:PARTICIPATES]->(userTournament:Tournament)-[:IS RELATED TO]->(game:Game)
    MATCH (game)<-[:IS RELATED TO]-(suggestedTournament:Tournament)
    WHERE NOT (user)-[:PARTICIPATES]->(suggestedTournament)
        AND suggestedTournament.startingTime > datetime()
        AND suggestedTournament.visibility = 'PUBLIC'
        AND NOT EXISTS {
            MATCH (suggestedTournament)<-[:PARTICIPATES]-(:User)
            WITH COUNT(*) AS participantCount
            WHERE participantCount >= suggestedTournament.maxParticipants
        }
    MATCH (user)-[:FOLLOWS]->(follower:User)-[:PARTICIPATES]->(suggestedTournament)
    WITH suggestedTournament, COUNT(DISTINCT follower) AS numParticipantFollowers
    MATCH (suggestedTournament)<-[:PARTICIPATES]-(participant:User)
    WITH suggestedTournament, numParticipantFollowers, COUNT(DISTINCT participant) AS numParticipants
    MATCH (suggestedTournament)-[:IS RELATED TO]->(relatedGame:Game)
    WITH suggestedTournament, numParticipantFollowers, numParticipants,
        relatedGame._id AS gameId,
        relatedGame.name AS gameName
    RETURN DISTINCT
        suggestedTournament._id AS id,
        suggestedTournament.name AS name,
        suggestedTournament.startingTime AS startingTime,
        suggestedTournament.maxParticipants AS maxParticipants,
        numParticipants,
        numParticipantFollowers,
        { id: gameId, name: gameName } AS game
    ORDER BY numParticipantFollowers DESC, suggestedTournament.startingTime ASC
    SKIP $pageSize * ($pageNumber - 1)
    LIMIT $pageSize
""")
List<TournamentSuggestionDTO> getTournamentsRecommendation(String username, int pageSize, int pageNumber);

```

## 9.3 Indexes

### 9.3.1 MongoDB indexes

This section explores potential indexes that could enhance the performance of read operations for each collection. Subsequently, we will evaluate each suggested index to determine whether its addition is justified.

#### Users Collection

- *username*: Adding an index on the *username* field is highly efficient because it optimizes the performance of operations like login, signup, and browsing users, which are among the most frequent actions performed on the database. Given the high frequency of these operations, creating an index on *username* provides a substantial

improvement in query performance, ensuring the system remains responsive and scalable as the database grows. The table illustrates the performance improvement achieved by using an index.

	<b>Without index</b>	<b>With index</b>
<b>Returned</b>	1	1
<b>Execution Time (ms)</b>	36	1
<b>Keys examined</b>	0	1
<b>Documents examined</b>	5000	1

## Games Collection

- *name*: An index has been added to the *name* field in the Games collection to significantly enhance the performance of the browse game operation, which is the primary activity of the application. This operation frequently involves searching for games by their names, filtering, or sorting them based on user input. With the index in place, the database can efficiently locate and retrieve games without performing a full collection scan. This not only reduces query execution time but also ensures scalability as the number of games increases. By prioritizing this optimization, the application can deliver faster response times and a smoother user experience during its most frequently used feature.

	<b>Without index</b>	<b>With index</b>
<b>Returned</b>	1	1
<b>Execution Time (ms)</b>	1	0
<b>Keys examined</b>	0	1
<b>Documents examined</b>	2000	1

## Reviews Collection

- *game.id*: Given the high frequency of queries that load reviews for a specific game or perform searches using *game.id*, we decided to add an index on the *game.id* field in the Reviews collection. This

index significantly enhances query performance by enabling faster and more efficient data retrieval. While maintaining the index adds some overhead during inserts, the benefits in read performance far outweigh the cost, especially in a read-heavy system. This ensures faster queries, better scalability, and an overall improved user experience. As demonstrated in the table below, adding the index significantly improves the performance of the Find reviews by *game.id* query, making it much more efficient and faster.

	<b>Without index</b>	<b>With index</b>
<b>Returned</b>	361	361
<b>Execution Time (ms)</b>	3321	6
<b>Keys examined</b>	0	361
<b>Documents examined</b>	674945	361

## Threads Collection and Tournament Collection

- *game.id*: As with the Reviews collection, we have also decided to add an index on *game.id* in the Threads and Tournaments collections for the same reasons mentioned earlier. This index significantly enhances query performance by enabling efficient retrieval of threads and tournaments related to specific games, avoiding costly full collection scans. While maintaining the index incurs some overhead during inserts or updates, the benefits in terms of faster read operations, improved scalability, and a smoother user experience far outweigh the costs.

	<b>Without index</b>	<b>With index</b>
<b>Returned</b>	50	50
<b>Execution Time (ms)</b>	11	0
<b>Keys examined</b>	0	50
<b>Documents examined</b>	9666	50

Table 9.1: Performance for Thread collection

	<b>Without index</b>	<b>With index</b>
<b>Returned</b>	20	20
<b>Execution Time (ms)</b>	449	2
<b>Keys examined</b>	0	20
<b>Documents examined</b>	20006	20

Table 9.2: Performance for Tournament collection

### 9.3.2 Neo4j indexes

#### User entity

- *username*: In the User entity on Neo4j, we deemed it necessary to add an index on the username field since all queries rely on the findByUsername operation. This index ensures efficient and fast retrieval of user data, significantly improving query performance and supporting the high frequency of operations involving username searches.

#### Game entity and Tournament entity

- *id*: In the Tournaments and Games entities, we added indexes on the id field because all Neo4j queries that start computations from a specific node use the id property to locate it. These indexes ensure efficient and fast access to nodes, significantly improving query performance and scalability.